


System on Chip (SoC) Design

Outline


- Key Trends and The SoC Paradigm
- System on Chip
 - Architecture
 - Design
 - Cores
 - Interconnection
- Cost Benefits of SoC
- Examples
- Conclusion

Moore's Law and Technology Scaling



...the performance of an IC, including the number components on it, doubles every 18-24 months with the same chip price ... - Gordon Moore - 1960

The Productivity Gap



- 100M logic gates in 90nm = Logic of 1000 ARM7's
- Current 0.13u SoC's: 10M\$ ~100M\$ design cost

ITRS Roadmap

Year of Technology Node	1999	2002	2005	2008	2011	2014
Technology	180nm	130nm	100nm	70nm	50nm	35nm
DRAM /introduction	1G	2~4G	8G	-	64G	-
Transistors/chip (μP) (M)	110	220~441	882	2,494	7,053	19,949
Chip size (μP) (mm^2)	450	450~567	622	713	817	937
Number of signal I/O (μP)	768	1,024	1,024	1,280	1,408	1,472
Power/Ground I/O (μP)	1,536	2,018	2,018	2,560	2,816	2,944
On-chip local clock (MHz) (high performance)	1,250	2,100	3,500	6,000	10,000	13,500
On-chip across-chip clock (MHz) (high performance)	1,200	1,600	2,000	2,500	3,000	3,600
Off-chip speed (MHz) (high perf., peripheral buses)	480	885	1,035	1,285	1,540	1,800
Power (W) H.P./H.H.	90/1.4	130/2.0	160/2.4	170/2.0	174/2.2	183/2.4
Power supply (V) H.P./H.H.	1.8/1.5	1.5/1.2	1.2/0.9	0.9/0.6	0.6/0.5	0.6/0.3
Metal levels # ($\mu\text{P}/\text{SoC}$)	7/6	8/7	9/8	9/9	10/10	10/10

H.P. – high performance microprocessor, μP – microprocessors
H.H. – hand-hold products, SoC – system-on-chip

Silicon technology roadmap

- intrinsic capability of ICs (transistor count / gate delay) grows with $\sim 50\%$ per year (Moore's Law)
- power limits the performance


	low power SoC			high performance MPU/SoC		
	2001	2004	2010	2001	2004	2010
gate length (nm)	130	90	45	90	53	25
supply voltage	1.2	1	0.6	1.1	1	0.6
transistor count (M)	3.3	8.3	40	276	553	2212
chip size (mm ²)	100	120	144	310	310	310
clock frequency (GHz)	0.15	0.3	0.6	1.7	2.4	4.7
wiring levels	6	7	9	7	8	10
max power (W)	0.1	0.1	0.1	130	160	218

Introduction - History

- First generation chips contained a few transistors.
- Today, silicon technology allows us to build chips consisting of hundreds of millions of transistors (Intel Pentium IV: 0.09 micron). This technology has enabled new levels of system integration onto a single chip.
- Mobile phones, portable computers and Internet appliances will be built using a single chip.
- The demand for more powerful products and the huge capacity of today's silicon technology have moved System-on-Chip (SoC) designs from leading edge to mainstream design practice.
- "System on Chip" (SoC) technology will put the maximum amount of technology into the smallest possible space.

Electronic systems

Systems on chip are everywhere



Technology advances enable increasingly more complex designs

Central Question: *how to exploit deep-submicron technologies efficiently?*

Main Challenges of Wireless Sensor Network

❑ Energy dissipation

- Reduce radiated power

- More power efficient radio

- Energy efficient protocols and routing algorithms

- Better trade-off between communication and local computing

❑ Size

- Higher integration (**System-on-Chip or SoC**)

❑ Cost


- Standard Digital CMOS Technology

Evolution of Microelectronics: the SoC Paradigm


Silicon Process Technology

- □ 0.13 μm CMOS
- □ ~100 millions of devices, 3 GHz internal Clock

Yesterday's chip is today's function block!



Paradigm Shift in SoC Design



System on a board



System on a Chip

Evolutionary Problems

Emerging new technologies:

- Greater complexity
- Increased performance
- Higher density
- Lower power dissipation

□ Key Challenges

- Improve productivity
- HW/SW codesign
- Integration of analog & RF IPs
- Improved DFT

□ Evolutionary techniques:

- IP (Intellectual Property) based design
- Platform-based design



Migration from ASICs to SoCs

ASICs are logic chips designed by end customers to perform a specific function for a desired application.

ASIC vendors supply libraries for each technology they provide. In most cases, these libraries contain predesigned and preverified logic circuits.

ASIC technologies are:


- ❖ gate array
- ❖ standard cell
- ❖ full custom

Migration from ASICs to SoCs

In the mid-1990s, ASIC technology evolved from a chip-set philosophy to an embedded-cores-based system-on-a-chip concept.

An SoC is an IC designed by stitching together multiple stand-alone VLSI designs to provide full functionality for an application.

An SoC compose of predesigned models of complex functions known as cores (terms such as intellectual property block, virtual components, and macros) that serve a variety of applications.



Three forms of SoC design

The scenario for SoC design is characterized by three forms:

1. **ASIC vendor design:** This refers to the design in which all the components in the chip are designed as well as fabricated by an ASIC vendor.
2. **Integrated design:** This refers to the design by an ASIC vendor in which all components are not designed by that vendor. It implies the use of cores obtained from some other source such as a core/IP vendor or a foundry.
3. **Desktop design:** This refers to the design by a fabless company that uses cores which for the most part have been obtained from other source such as IP companies, EDA companies, design services companies, or a foundry.

SoC Design Challenges

Why does it take longer to design SOC's compared to traditional ASIC's?

We must examine factors influencing the degree of difficulty and **Turn Around Time (TAT)** (the time taken from gate-level netlist to metal mask-ready stage) for designing ASIC's and SOC's.



For an ASIC, the following factors influence TAT:

- Frequency of the design
- Number of clock domains
- Number of gates
- Density
- Number of blocks and sub-blocks

The key factor that influences TAT for SOC's is **system integration** (integrating different silicon IPs on the same IC).

SoC Design Challenges

Levarage Internal Bandwidth vs External Bandwidth



SoCs vs. ASICs

- ❑ SoC is not just a large ASIC
 - ❖ Architectural approach involving significant design reuse
 - ❖ Addresses the cost and time-to-market problems
- ❑ SoC methodology is an incremental step over ASIC methodology
- ❑ SoC design is significantly more complex
 - ❖ Need cross-domain optimizations
 - ❖ IP reuse and Platform-based design increase productivity, but not enough
 - ❖ Even with extensive IP reuse, many of the ASICs design problems remain, plus many more ...
 - ❖ Productivity increase far from closing design gap

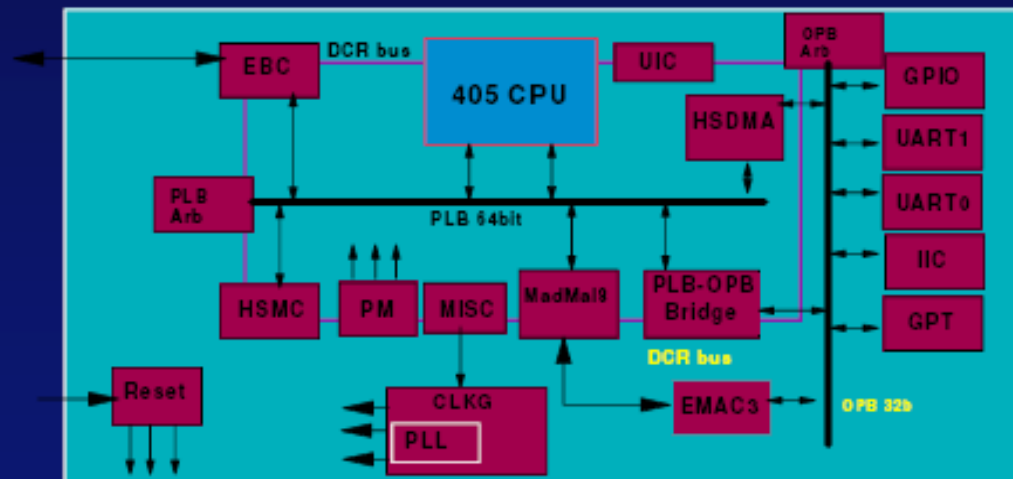
From ASICs to SoCs

Much higher
levels of
integration
($< .25\mu\text{m}$)

High
Performance
Circuits and
Tools


Common
Architectural
Blocks

Shorter design
cycle
+ Lagging
Productivity



- IP Cores
- Platforms
- Reuse

Technology vs. Productivity vs. Complexity



System on Chip benefits

Typical approach :

Define requirements

Design with off-the shelf chips


- at 0.5 year mark : first prototypes
- 1 year : ship with low margins/loss
start ASIC integration
- 2 years : ASIC-based prototypes
- 2.5 years : ship, make profits (with competition)

With SoC

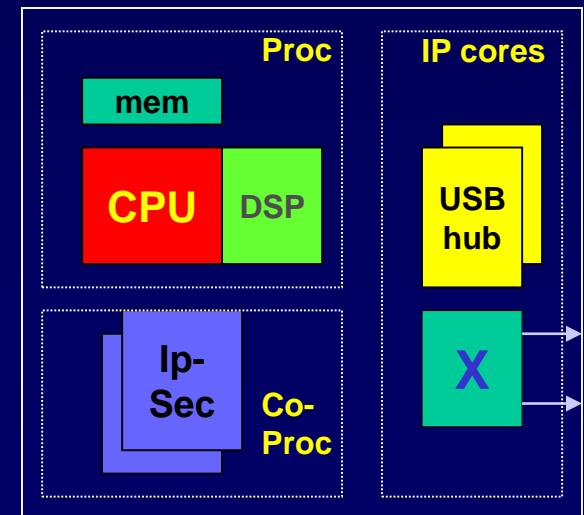
Define requirements

Design with off-the shelf cores

- at 0.5 year mark : first prototypes
- 1 year : ship with high margin and market share



Now : collection of cores



Typical : \$10

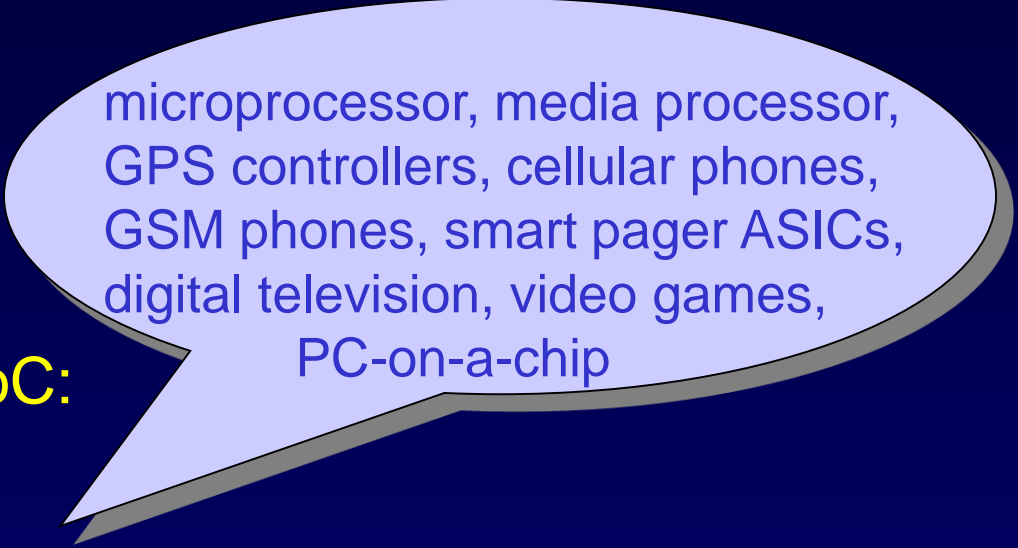
Typical applications of SoC

An SoC is a system on an IC that integrates software and hardware Intellectual Property (IP) using more than one design methodology for the purpose of defining the functionality and behavior of the proposed system.

The designed system is application specific.

Typical applications of SoC:

- ❖ consumer devices,
- ❖ networking,
- ❖ communications, and
- ❖ other segments of the electronics industry.



microprocessor, media processor,
GPS controllers, cellular phones,
GSM phones, smart pager ASICs,
digital television, video games,
PC-on-a-chip

A common set of problems facing everyone who is designing complex chips

- Time-to-market pressures demand rapid development.
- Quality of results (performance, area, power) - key to market success.
- Increasing chip complexity makes verification more difficult.
- Deep submicron issues make timing closure more difficult.
- The development team has different levels and areas of expertise, and is often scattered throughout the world.
- Design team members may have worked on similar designs in the past, but cannot reuse these designs because the design flow, tools, and guidelines have changed.
- SoC designs include embedded processor cores, and thus a significant software component, which leads to additional methodology, process, and organizational challenges.

Reusing macros (called “cores”, or IP) that have already been designed and verified helps to address all of the problems above.

Design for Reuse

To overcome the design gap, design reuse - the use of pre-designed and pre-verified cores, or reuse of the existing designs becomes a vital concept in design methodology.

An effective block-based design methodology requires an extensive library of reusable blocks, or macros, and it is based on the following principles:

- ❖ The macro must be extremely easy to integrate into the overall chip design.
- ❖ The macro must be so robust that the integrator has to perform essentially no functional verification of internals of the macro.

The challenge for designers is not whether to adopt reuse, but how to employ it effectively.

Design for Reuse

To be fully reusable, the hardware **macro** must be:

- **Designed to solve a general problem** – easily configurable to fit different applications.
- **Designed for use in multiple technologies** – For soft macros, this mean that the synthesis scripts must produce satisfactory quality of results with a variety of libraries. For hard macros, this means having an effective porting strategy for mapping the macro onto new technologies.
- **Designed for simulation with a variety of simulators** – Good design reuse practices dictate that both a Verilog and VHDL version of each model and verification testbench should be available, and they should work with all the major commercial simulators.
- **Designed with standards-based interfaces** – Unique or custom interfaces should be used only if no standards-based interface exists.

Design for Reuse - cont.

To be fully reusable, the hardware **macro** must be:


- **Verified independently of the chip in which it will be used** – Often, macros are designed and only partially tested before being integrated into a chip for verification. Reusable designs must have full, stand-alone testbenches and verification suites that afford very high levels of test coverage.
- **Verified to a high level of confidence** – This usually means very rigorous verification as well as building a physical prototype that is tested in an actual system running real software.
- **Fully documented in terms of appropriate applications and restrictions** – In particular, valid configurations and parameter values must be documented. Any restrictions on configurations or parameter values must be clearly stated. Interfacing requirements and restrictions on how the macro can be used must be documented.

Intellectual Property

Utilizing the predesigned modules enables:

- ❖ to avoid reinventing the wheel for every new product,
- ❖ to accelerate the development of new products,
- ❖ to assemble various blocks of a large ASIC/SoC quite rapidly,
- ❖ to reduce the possibility of failure based on design and verification of a block for the first time.

Resources vs. Number of Uses



These predesigned modules are commonly called Intellectual Property (IP) cores or Virtual Components (VC).

Intellectual Property Categories

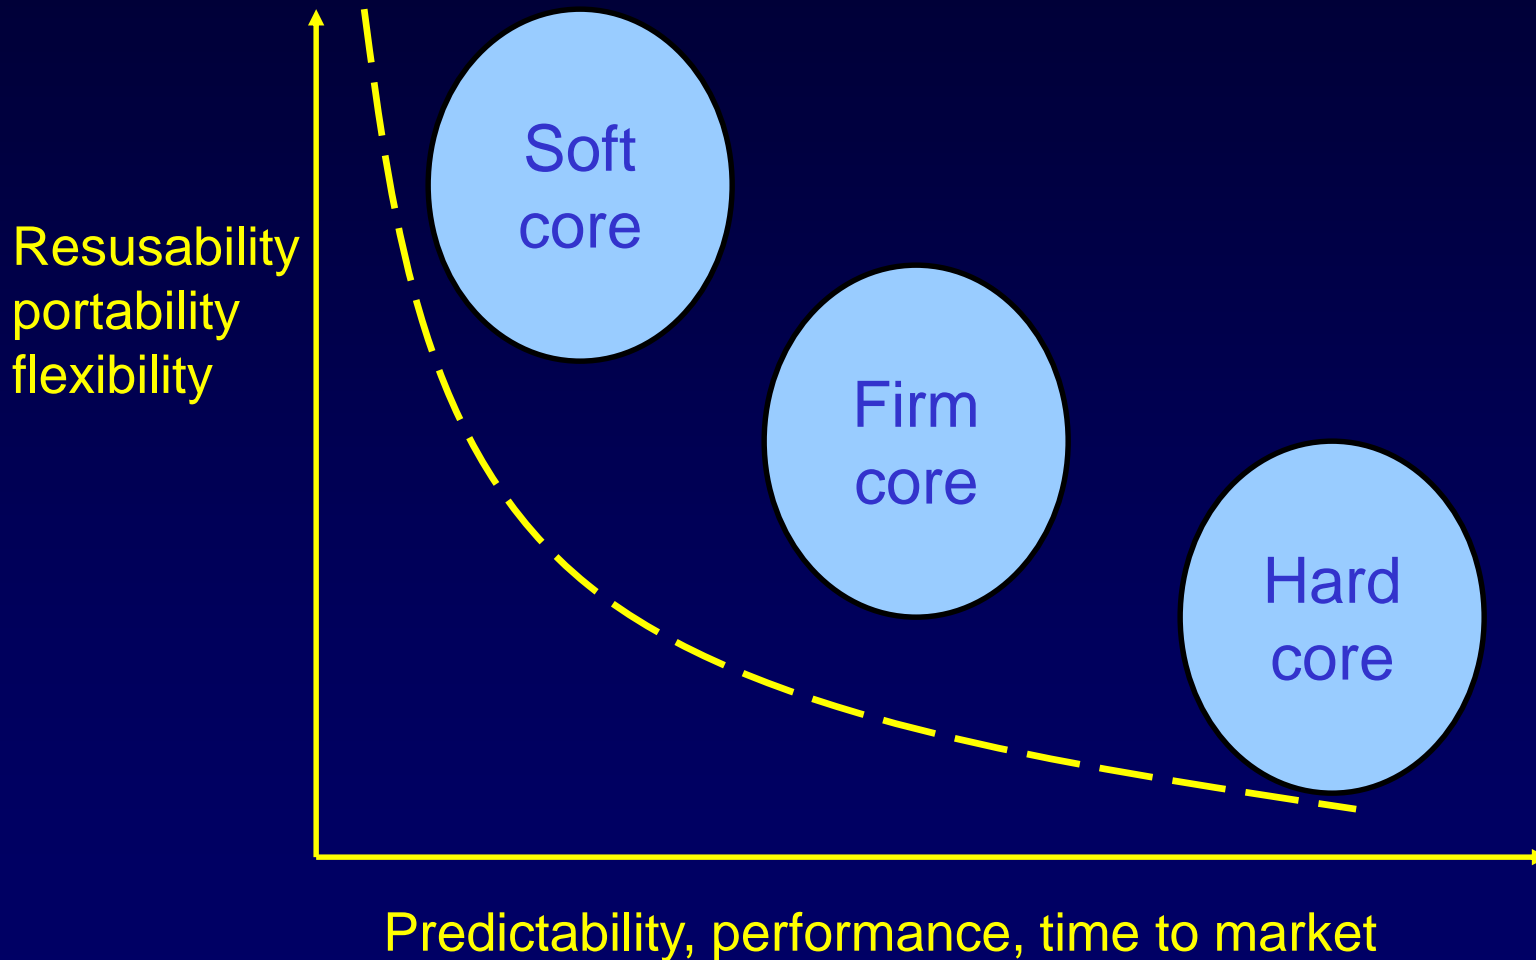
IP cores are classified into three distinct categories:

Hard IP cores consist of hard layouts using particular physical design libraries and are delivered in masked-level designed blocks (GDSII format). The integration of hard IP cores is quite simple, but hard cores are technology dependent and provide minimum flexibility and portability in reconfiguration and integration.

Soft IP cores are delivered as RTL VHDL/Verilog code to provide functional descriptions of IPs. These cores offer maximum flexibility and reconfigurability to match the requirements of a specific design application, but they must be synthesized, optimized, and verified by their user before integration into designs.

Firm IP cores bring the best of both worlds and balance the high performance and optimization properties of hard IPs with the flexibility of soft IPs. These cores are delivered in form of targeted netlists to specific physical libraries after going through synthesis without performing the physical layout.

Trade-offs among soft, firm, and hard cores




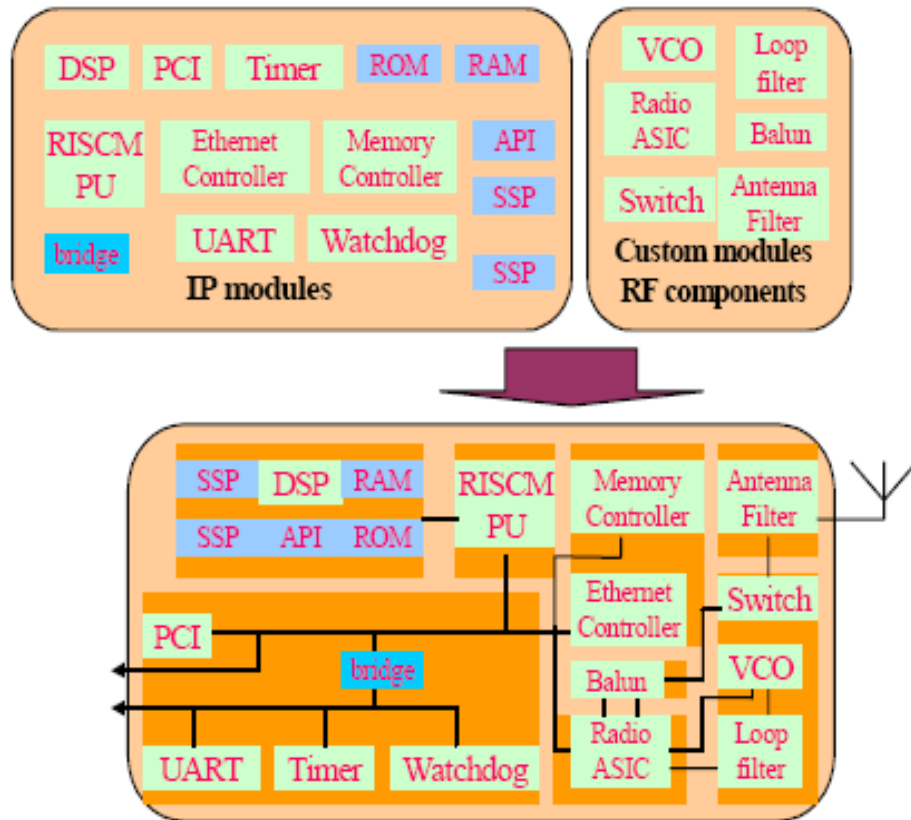
Comparison of Different IP Formats

IP Format	Representation	Optimization	Technology	Reusability
Hard	GDSII	Very High	Technology Dependent	Low
Soft	RTL	Low	Technology Independent	Very High
Firm	Target Netlist	High	Technology Generic	High


Examples of IPs

Category	Intellectual Property
Processor	ARM7, ARM9, and ARM10, ARC
Application-Specific DSP	ADPCM, CELP, MPEG-2, MPEG-4, Turbo Code, Viterbi, Reed Solomon, AES
Mixed Signal	ADCs, DACs, Audio Codecs, PLLs, OpAmps, Analog MUX
I/Os	PCI, USB, 1394, 1284, E-IDE, IRDA
Miscellaneous	UARTs, DRAM Controller, Timers, Interrupt Controller, DMA Controller, SDRAM Controller, Flash Controller, Ethernet 10/100 MAC

IP Reuse and IP-Based SoC Design



Yesterday/Today: Real Components



Bluetooth 

Today/Future: Virtual Components

What is MPSoC

MPSoC is a system-on-chip that contains multiple instruction-set processors (CPUs).

The typical MPSoC is a **heterogeneous multiprocessor**: there may be several different types of processing elements (PEs), the memory system may be heterogeneously distributed around the machine, and the interconnection network between the PEs and the memory may also be heterogeneous.

MPSoCs often require large amounts of memory. The device may have embedded memory on-chip as well as relying on off-chip commodity memory.

The design process of SoCs

SoC designs are made possible by deep submicron technology. This technology presents a whole set of design challenges including:

- ❖ **interconnect delays,**
- ❖ **clock and power distribution, and**
- ❖ **the placement and routing of millions of gates.**

These physical design problems can have a significant impact on the functional design of SoCs and on the design process itself.

The first step in system design is specifying the required functionality. The second step is to transform the system functionality into an architecture which define the system implementation by specifying the number and types of components and connections between them.

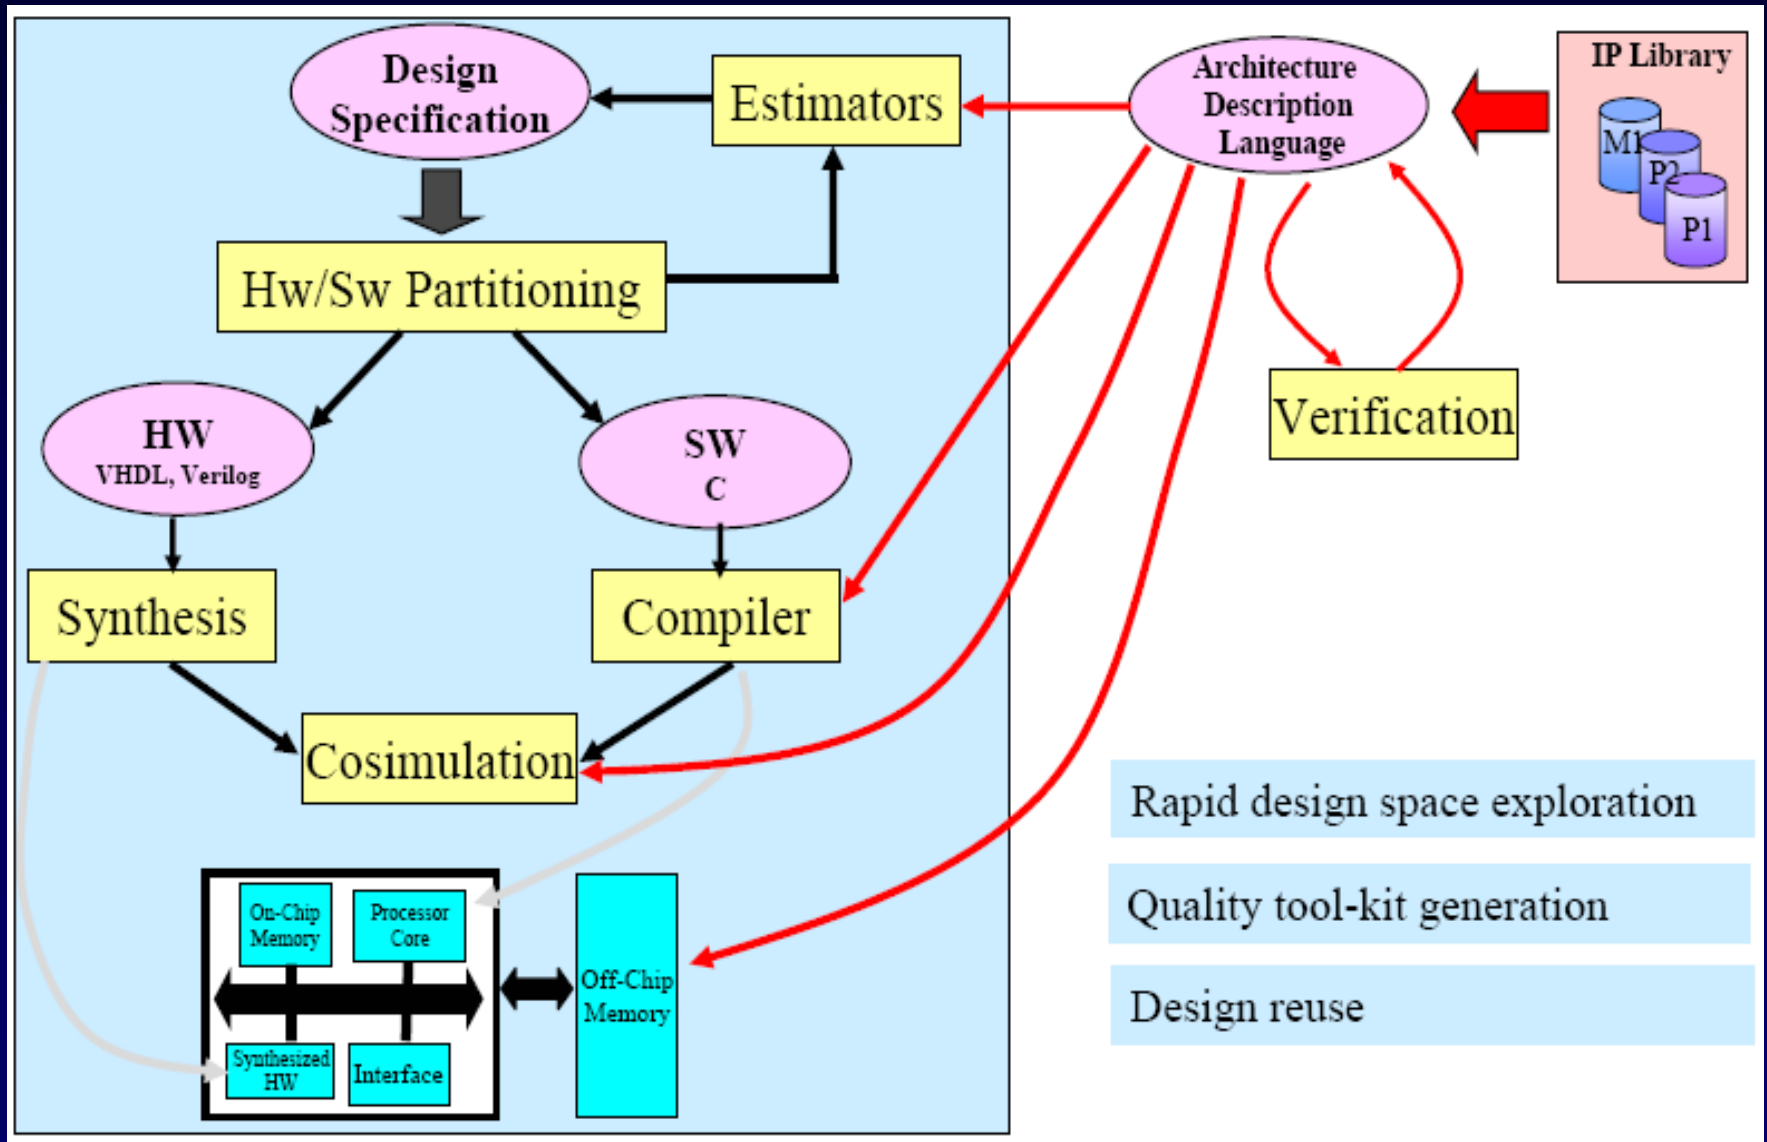
Define Hardware-Software Codesign

Hardware-Software Codesign is the concurrent and co-operative design of hardware and software components of a system.

The SoC design process is a hardware-software codesign in which design productivity is achieved by design reuse.

The design process is the set of design tasks that transform an abstract specification model into an architectural model.

SoC Co-design Flow

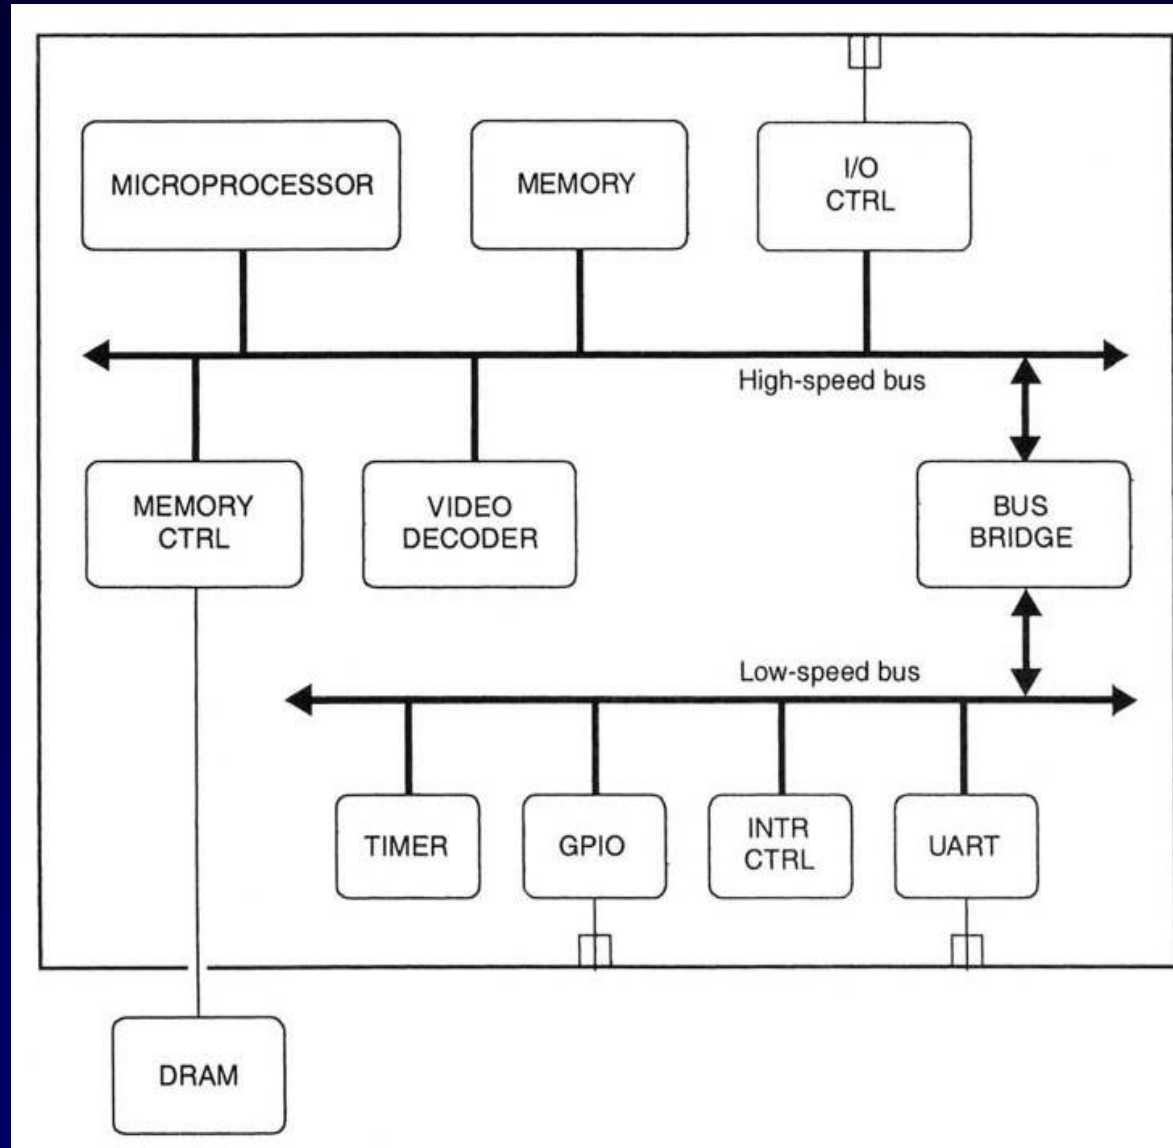


Design Proces

A canonical or generic form of an SoC design

These chips have:

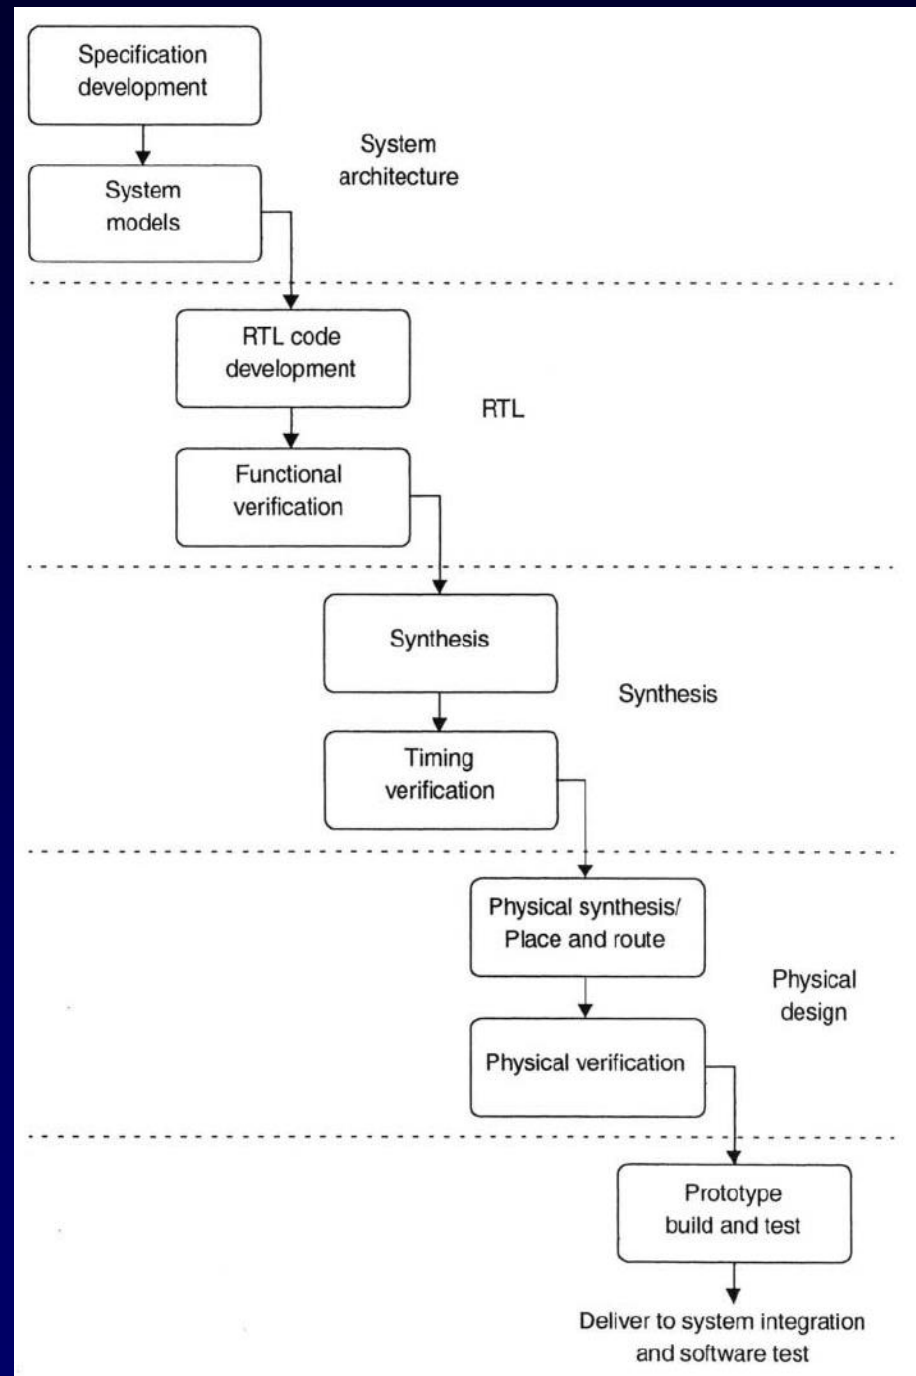
- one (several) processors
- large amounts of memory
- bus-based architectures
- peripherals
- coprocessors
- and I/O channels



Waterfall vs. Spiral Design Flow

The traditional model for ASIC development is often called a **waterfall model**.

The project transitions from phase to phase in a step function, never returning to the activities of the previous phase.



Waterfall vs. Spiral Design Flow

As complexity increases, geometry shrinks, and time-to-market pressures continue to escalate, chip designers are moving from the old waterfall model to the newer **spiral development model**.

In the spiral model, the design team works on multiple aspects of the design simultaneously, incrementally improving in each area as the design converges on completion.

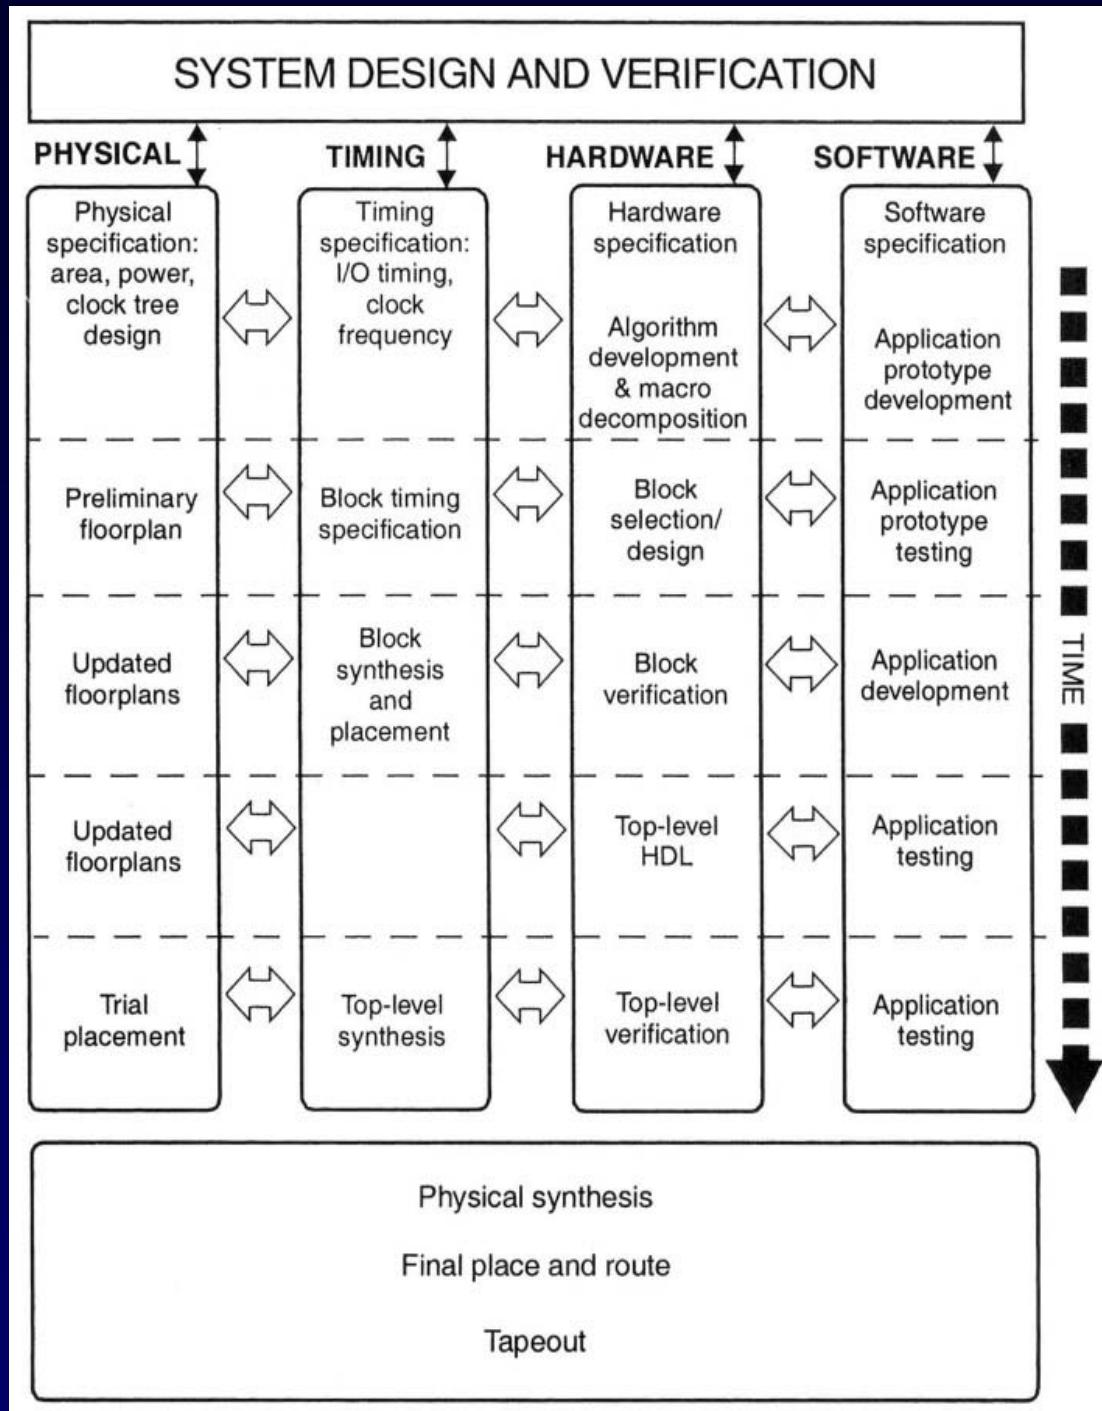
The spiral SoC design flow is characterized by:

- ❖ Parallel, concurrent development of hardware and software
- ❖ Parallel verification and synthesis of modules
- ❖ Floorplanning and place-and-route included in the synthesis process
- ❖ Modules developed only if a pre-designed hard or soft macro is not available
- ❖ Planned iteration throughout

Waterfall vs. Spiral Design Flow

Spiral SoC Design Flow

Goal: Maintain parallel interacting design flow



Top-Down vs. Bottom-Up

The classic top-down design process can be viewed as a recursive routine that begins with specification and decomposition, and ends with integration and verification:

- ❖ Write complete specifications for the system or subsystem being designed.
- ❖ Refine its architecture and algorithms, including software design and hardware/software cosimulation if necessary.
- ❖ Decompose the architecture into well-defined macros.
- ❖ Design or select macros; this is where the recursion occurs.
- ❖ Integrate macros into the top level; verify functionality and timing.
- ❖ Deliver the subsystem/system to the next higher level of integration; at the top level, this is tapeout.
- ❖ Verify all aspects of the design (functionality, timing, etc.).

Top-Down vs. Bottom-Up

A top-down methodology assumes that the lowest level blocks specified can, in fact, be designed and built. If it turns out that a block is not feasible to design, the whole specification process has to be repeated.

For this reason, real world design teams usually use a mixture of top-down and bottom-up methodologies, building critical low-level blocks while they refine the system and block specifications.

Libraries of reusable hard and soft macros clearly facilitate this process by providing a source of pre-verified blocks, proving that at least some parts of the design can be designed and fabricated in the target technology and perform to specification.

Design processes in flow diagrams

The first part of the design process consists of recursively developing, verifying, and refining **a set of specifications** until they are detailed enough to allow RTL coding to begin.

The specifications must completely describe all the interfaces between the design and its environment, including:

- ❖ **Hardware** – Functionality; External interfaces to other hardware (pins, buses, and how to use them); Interface to SW (register definitions); Timing; Performance; Physical design issues such as area and power
- ❖ **Software** – Functionality; Timing; Performance; Interface to HW SW structure, kernel

Type of specifications:

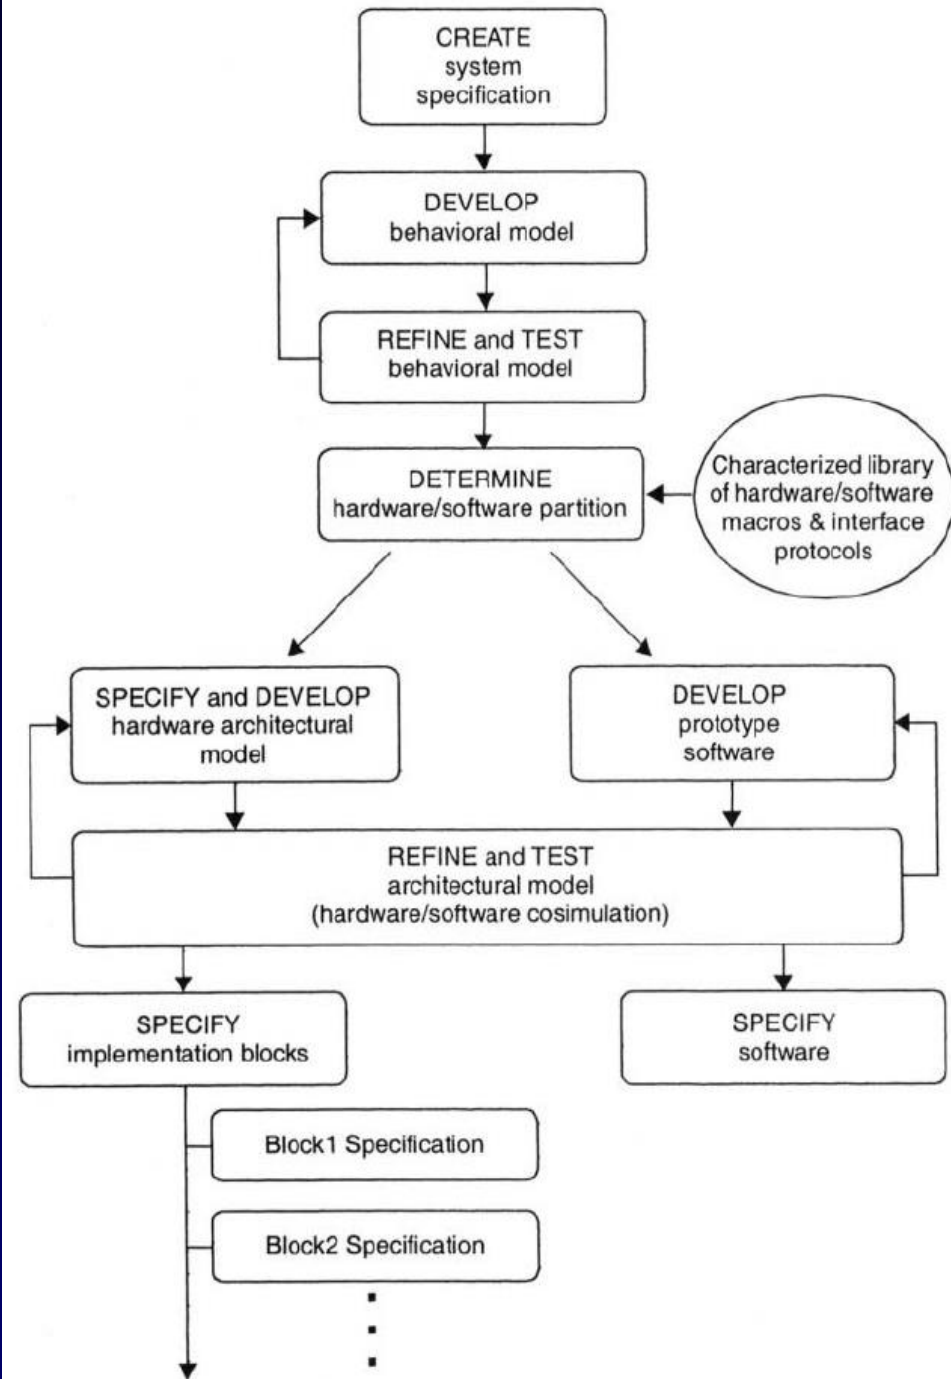
- ❖ **Formal specifications** – the desired characteristics of the design are defined independently of any implementation.
- ❖ **Executable specifications** – are typically an abstract model for the hardware and/or software being specified, and currently more useful for describing functional behavior in most design situations.

The System Design Process

Determining the optimal architecture (cost and performance) involves a set of complex decisions, such as:

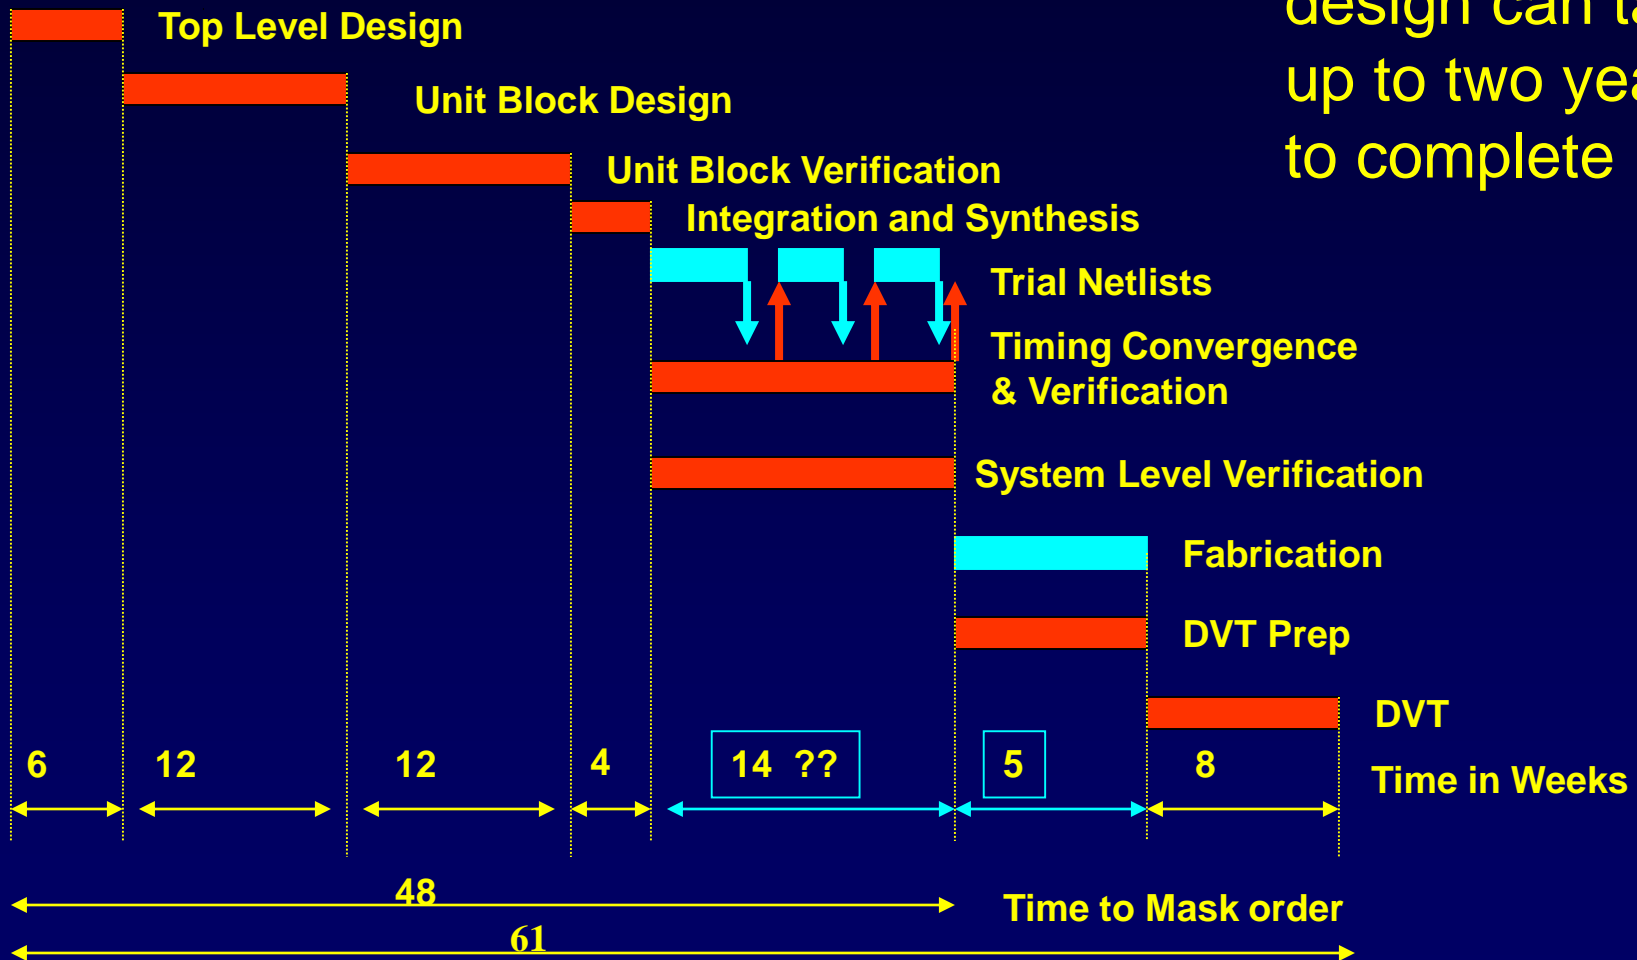
- What goes in software and what goes in hardware
- What processor(s) to use, and how many
- What bus architecture is required to achieve the required system performance
- What memory architecture to use to reach an appropriate balance between power, area, and speed.

Solution: modeling of several alternative architectures

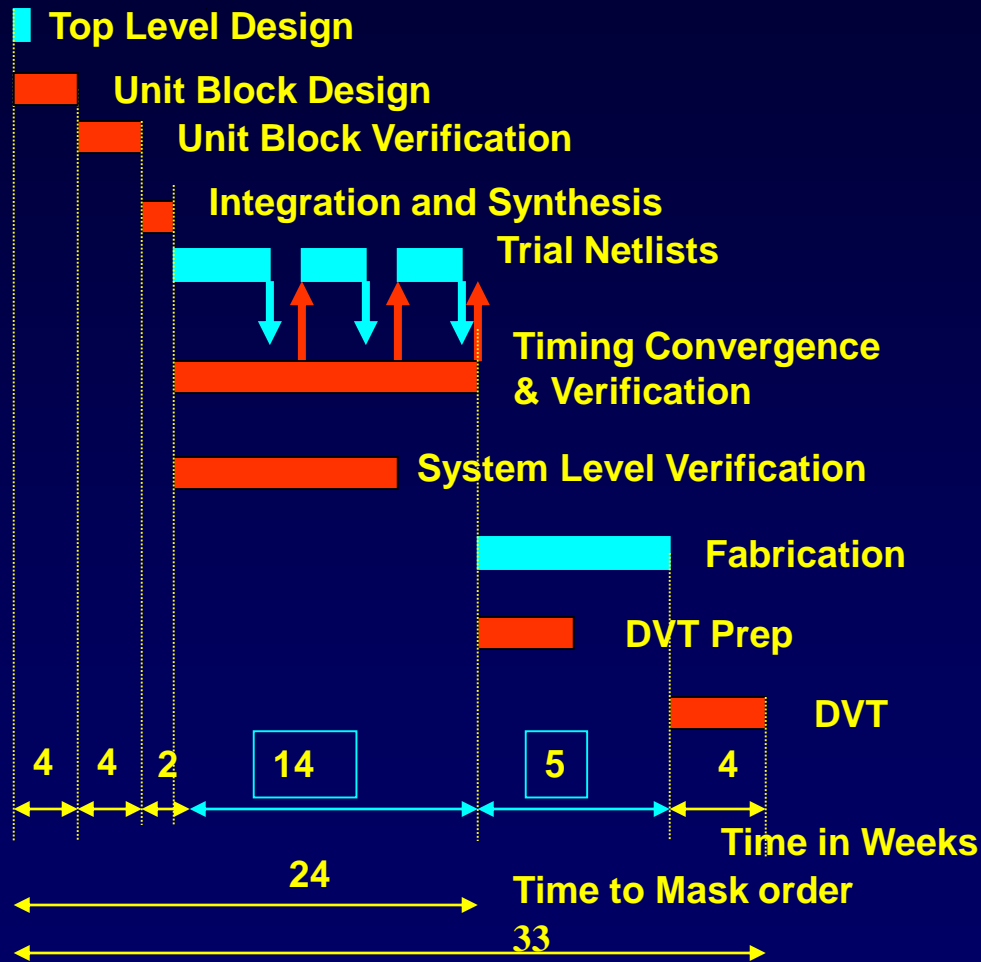


ASIC Typical Design Steps

Typical ASIC design can take up to two years to complete



SoC Typical Design Steps



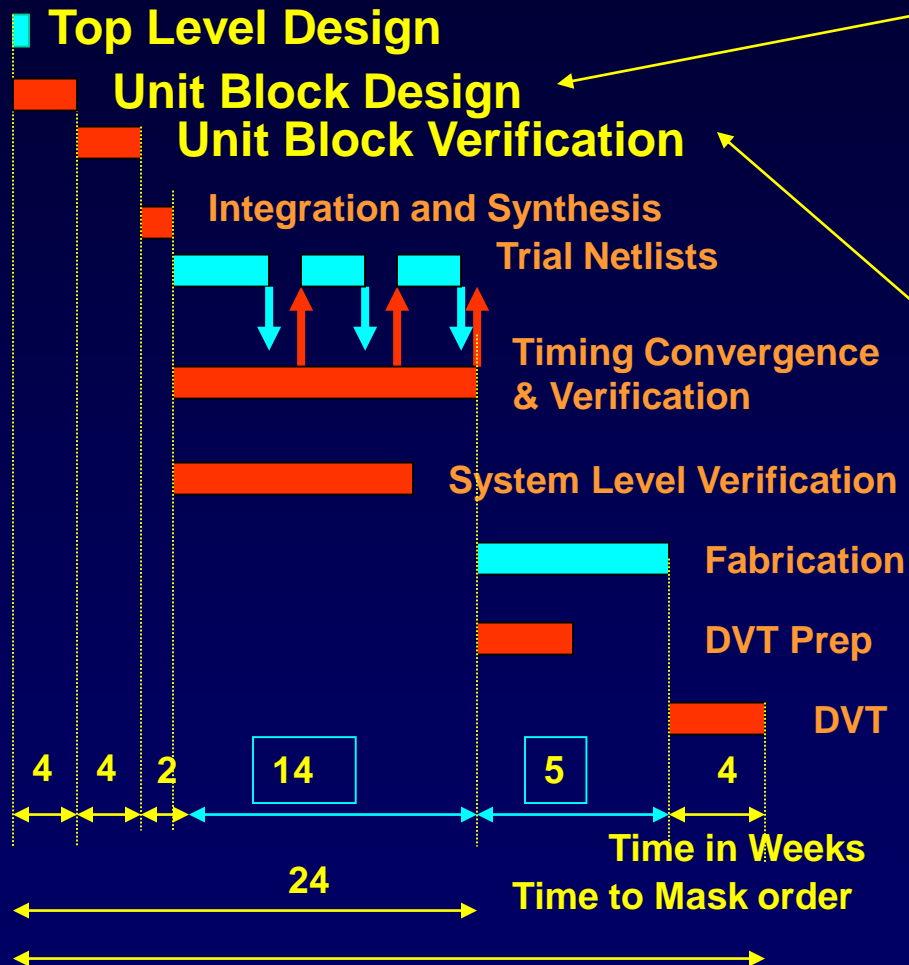
- With increasing Complexity of IC's and decreasing Geometry, IC Vendor steps of Placement, Layout and Fabrication are unlikely to be greatly reduced.

- In fact there is a greater risk that Timing Convergence steps will involve more iteration.

- Need to reduce time before Vendor Steps.

- Need to consider Layout issues up-front.

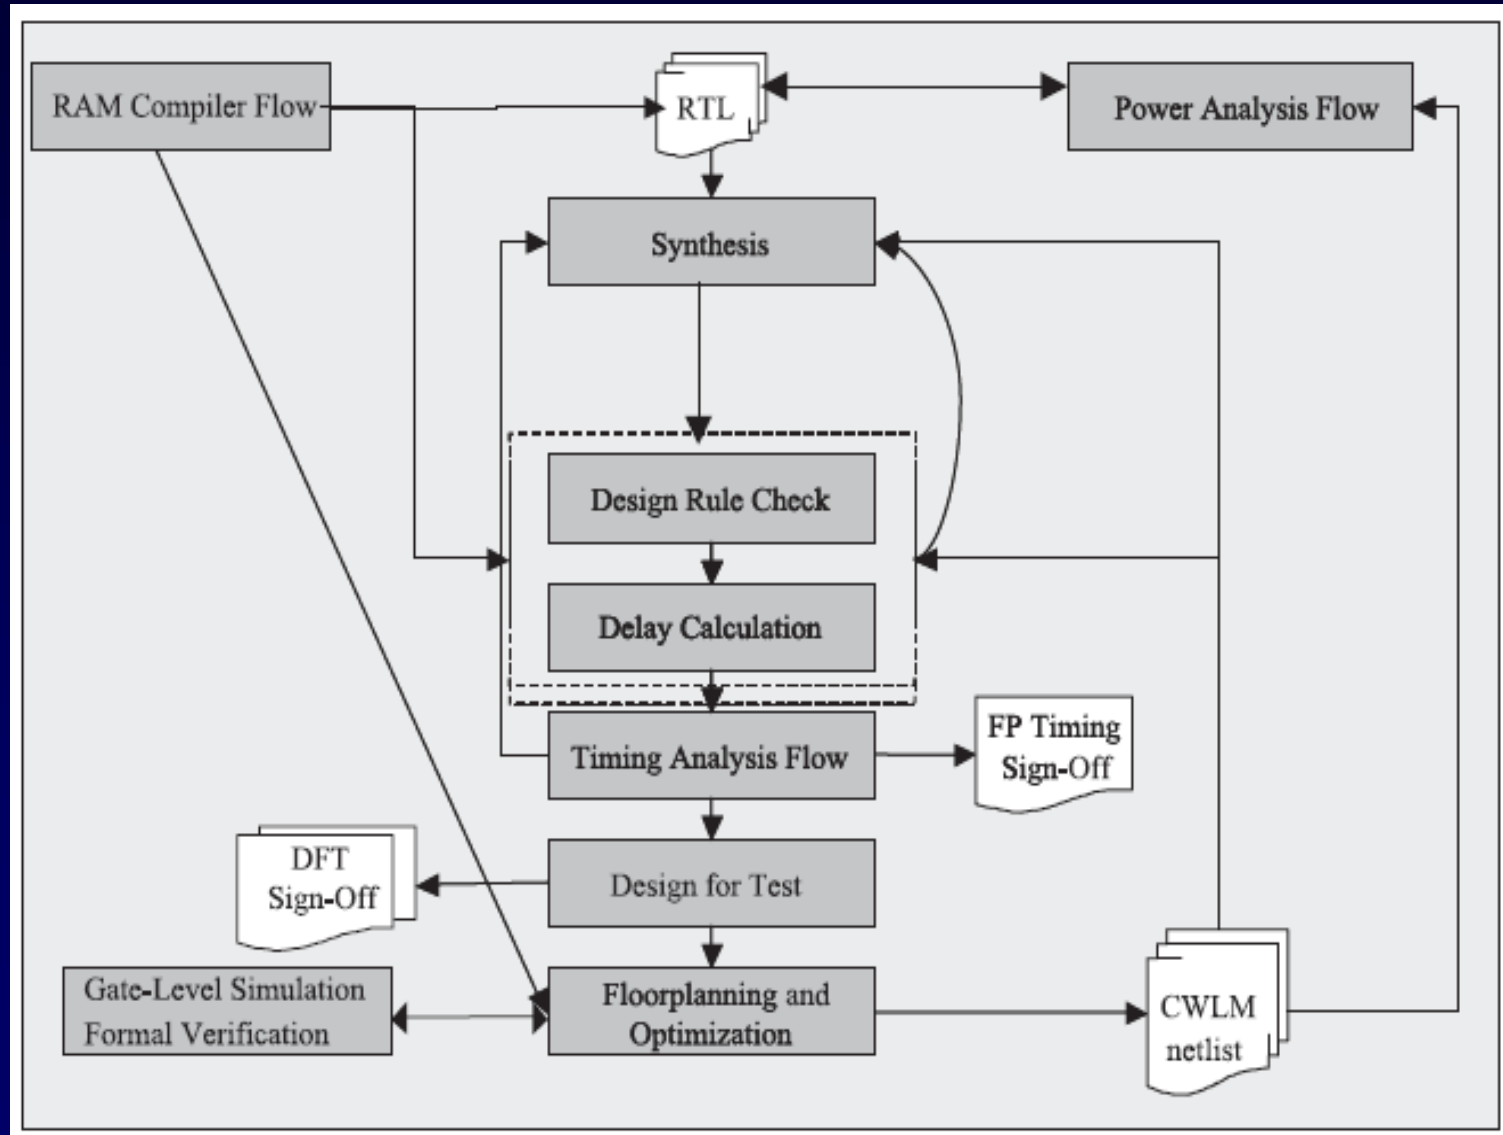
SoC Typical Design Steps



- SoC Architecture already defined. Flexible to scale in frequency and complexity. Allows new IP cores, new technology to be integrated.
- Separate the design of the reusable IP from the design of the SoC. Build the SoC from library of tested IP.
- Unit design consists only of any additional core features or wrapping new IP to enable integration.
- Reusable IP purchased from external sources, developed from in-house designs or designed as separate project off critical SoC development path.

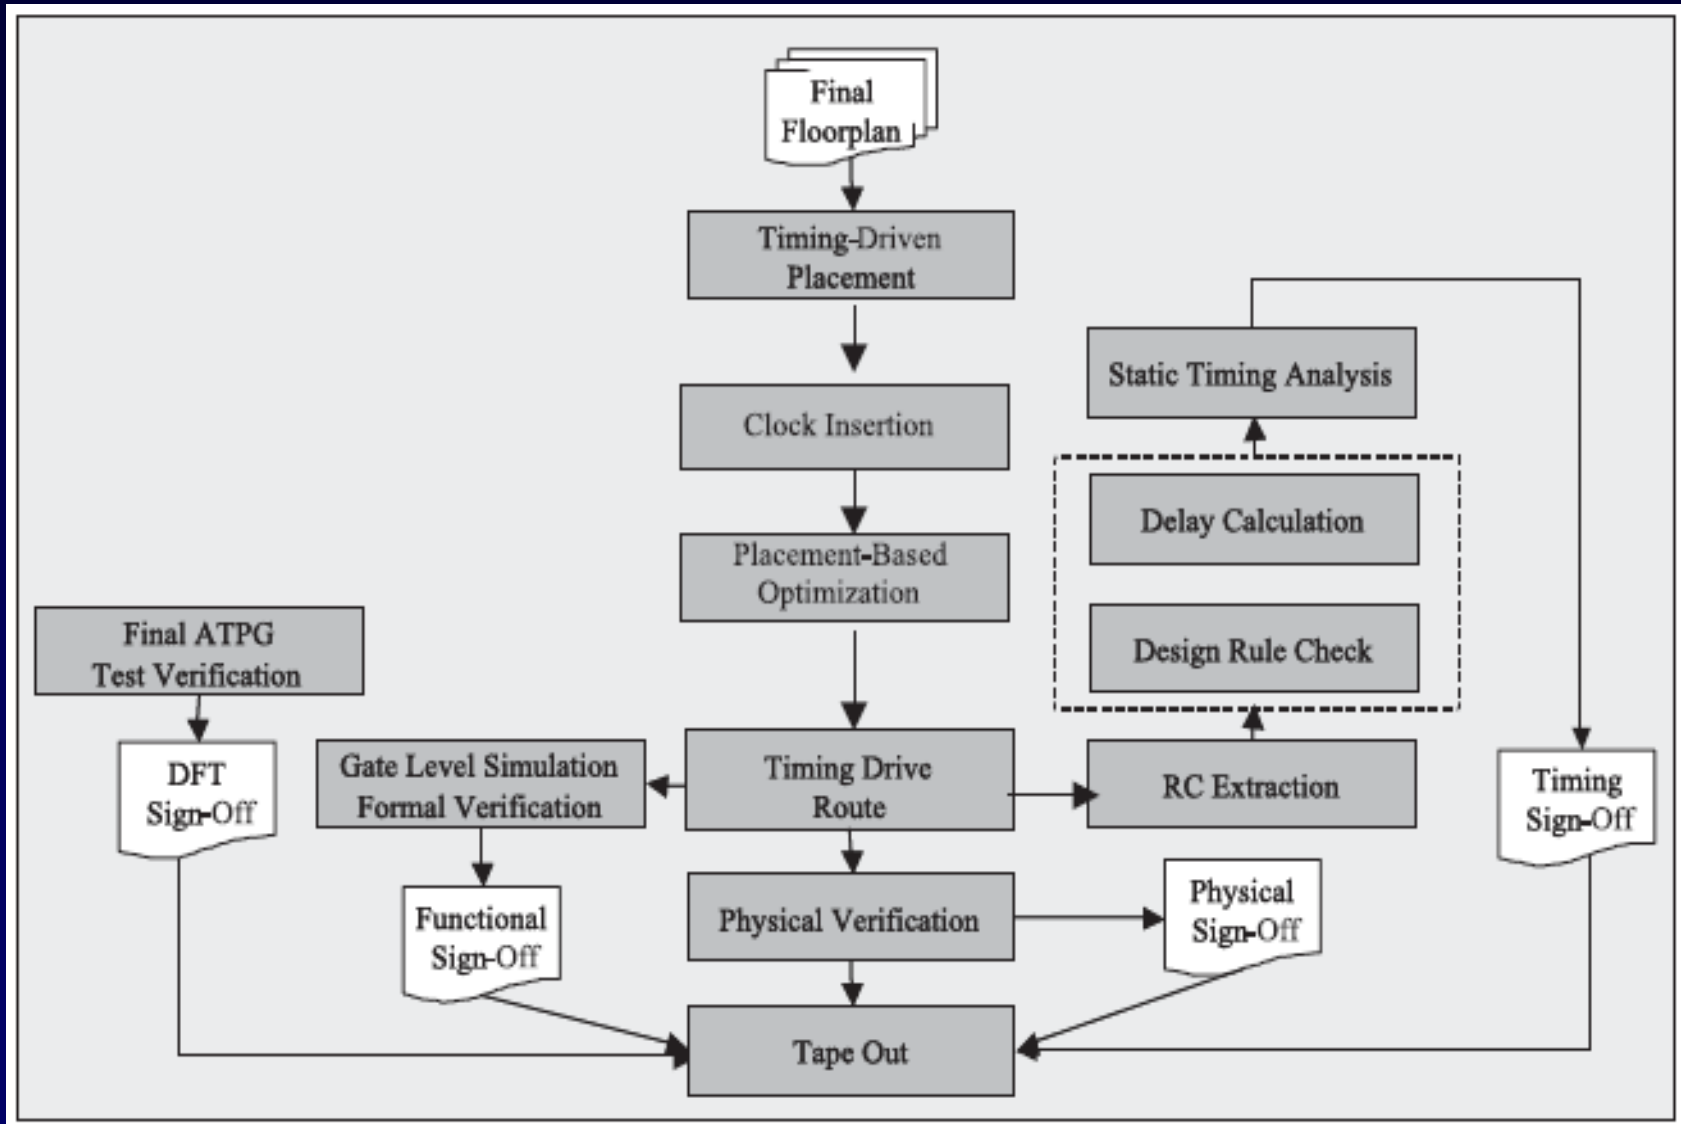
Design Methodology

A Front-End ASIC Design Flow

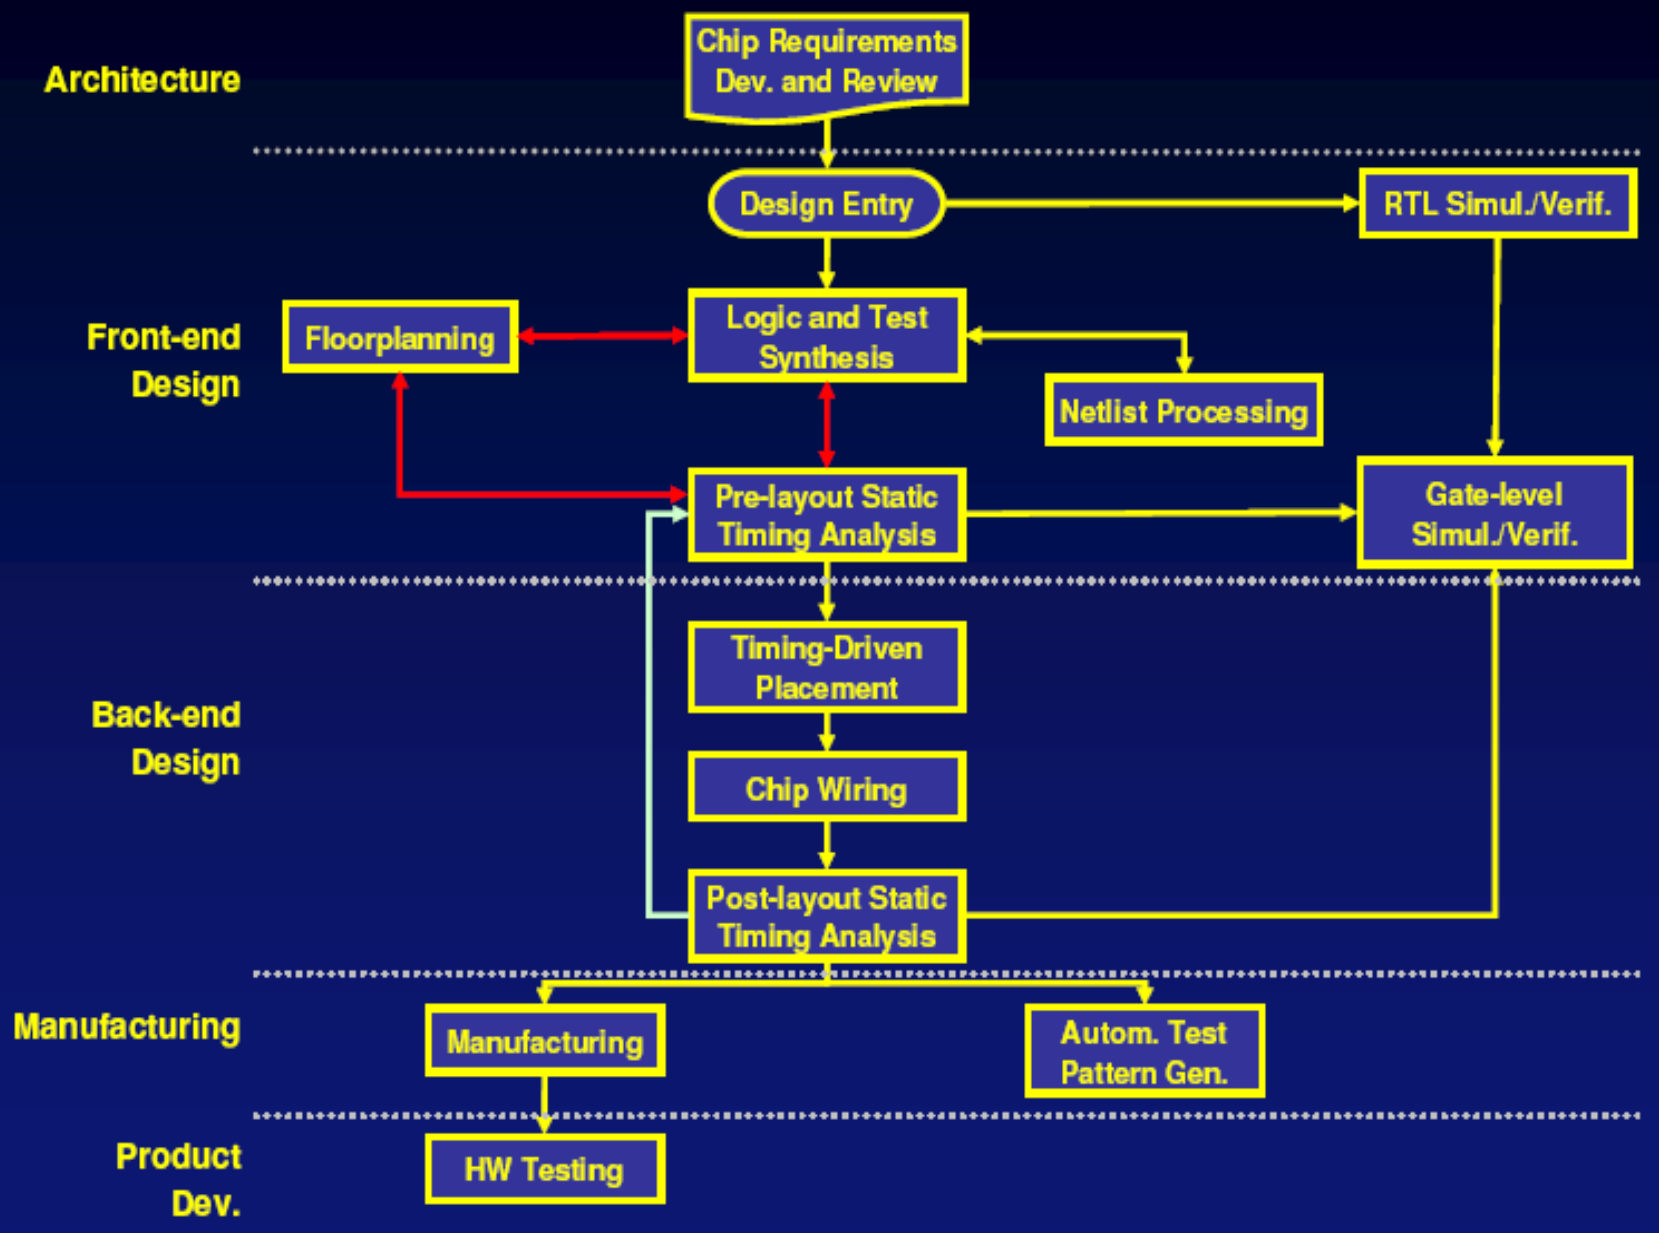


Design Methodology

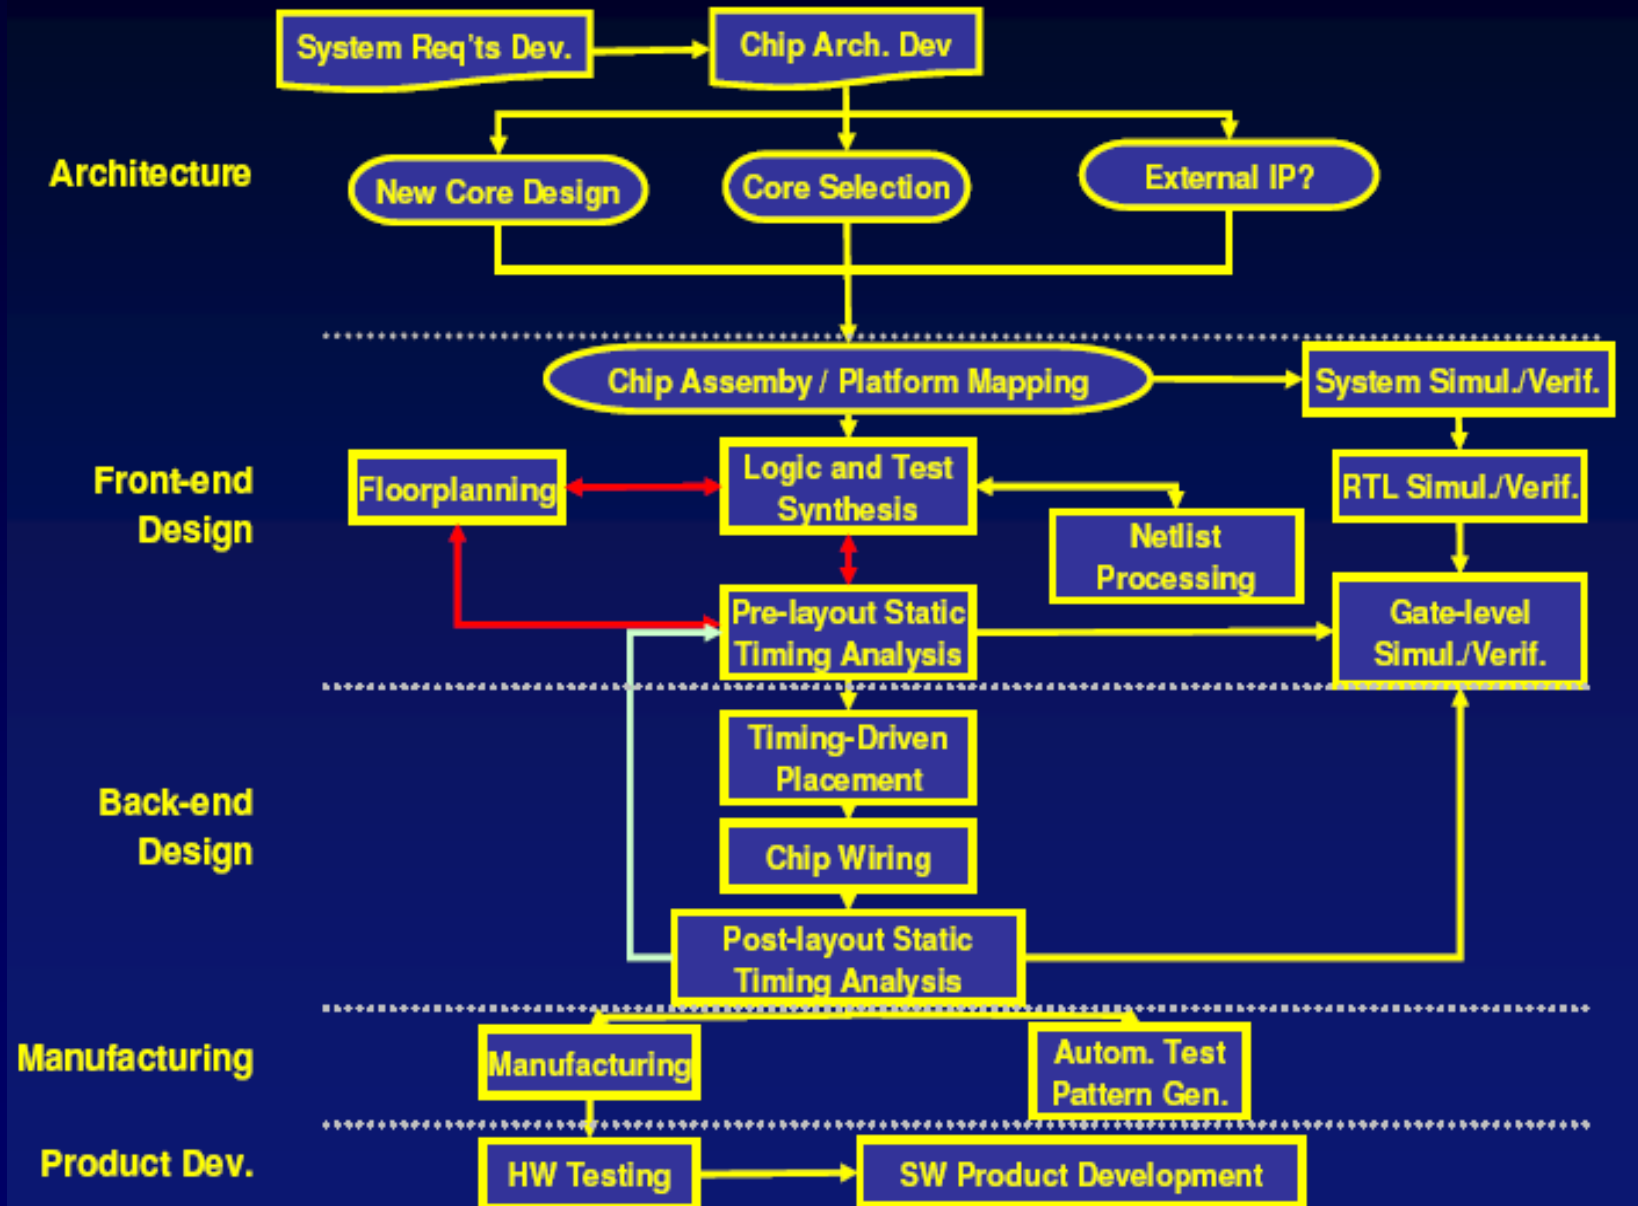
A Back-End Design Flow or Generic Physical Flow.



ASIC Methodology

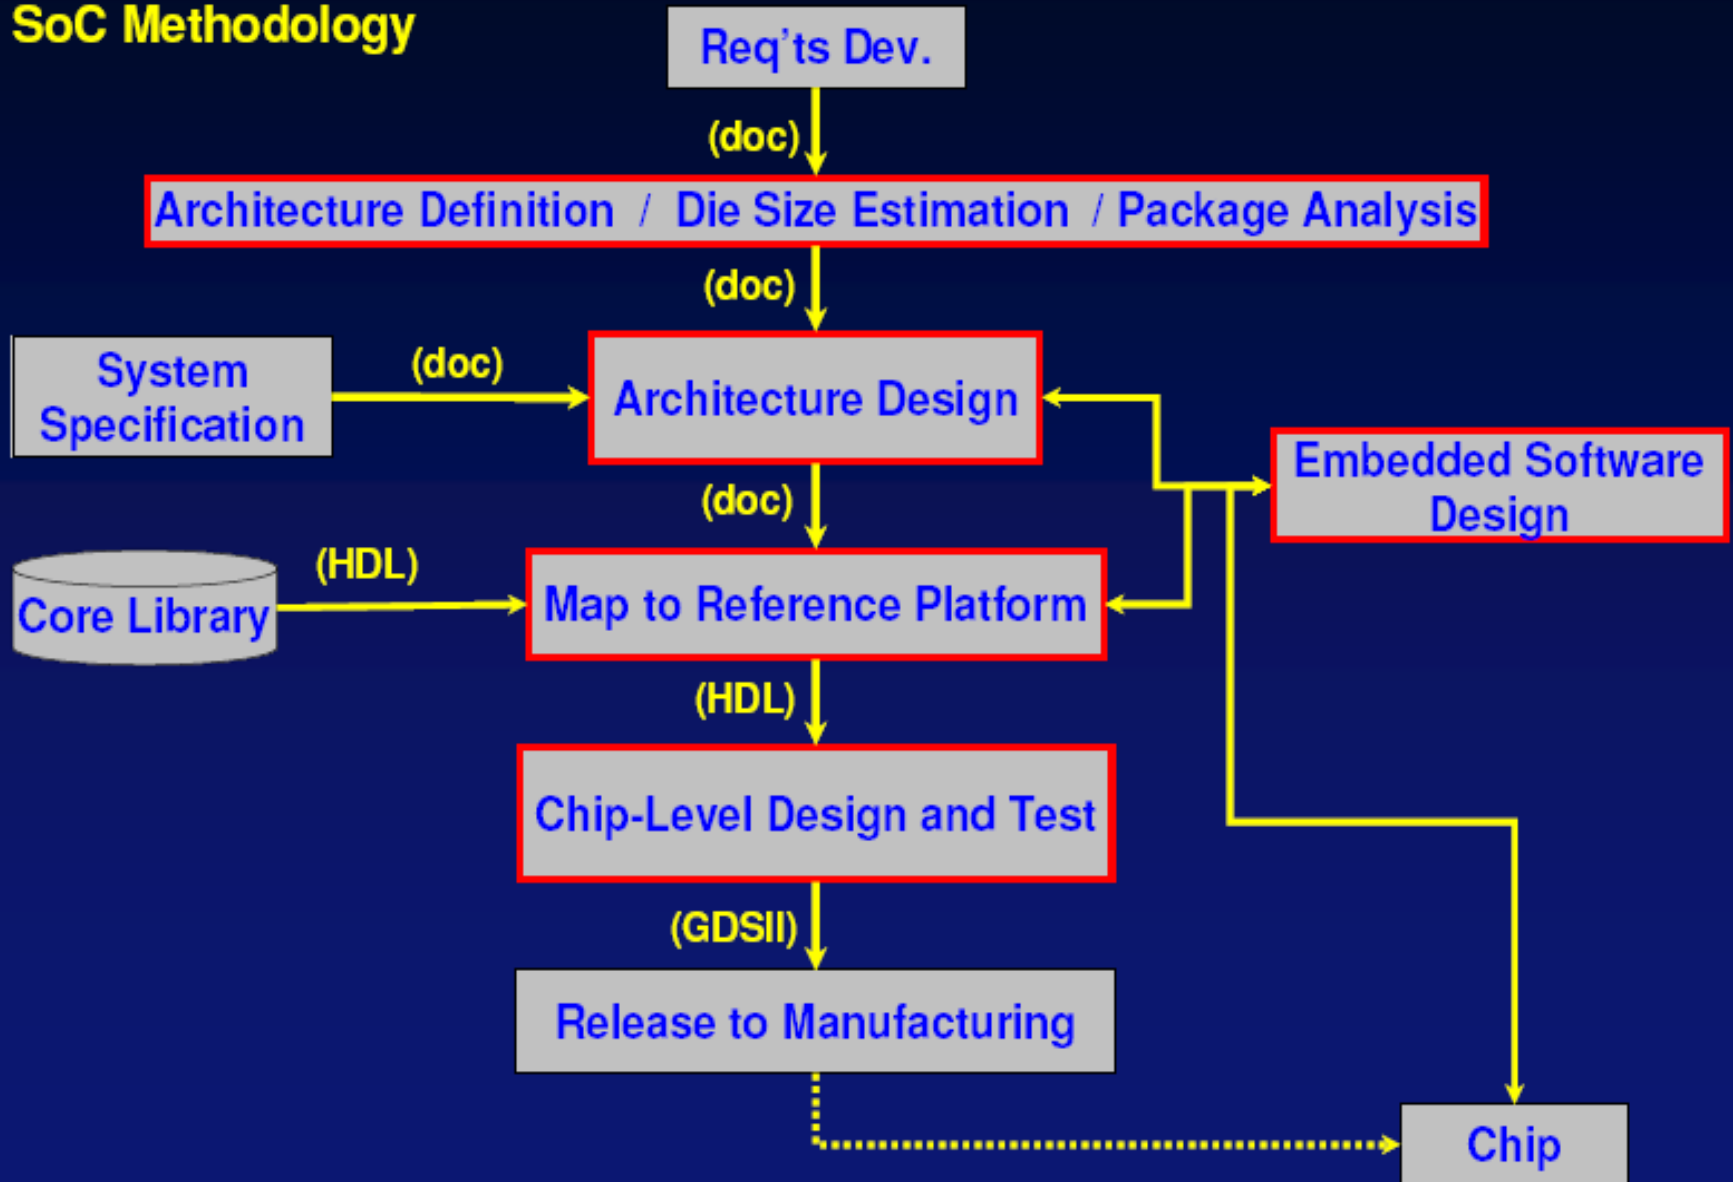


SOC Methodology



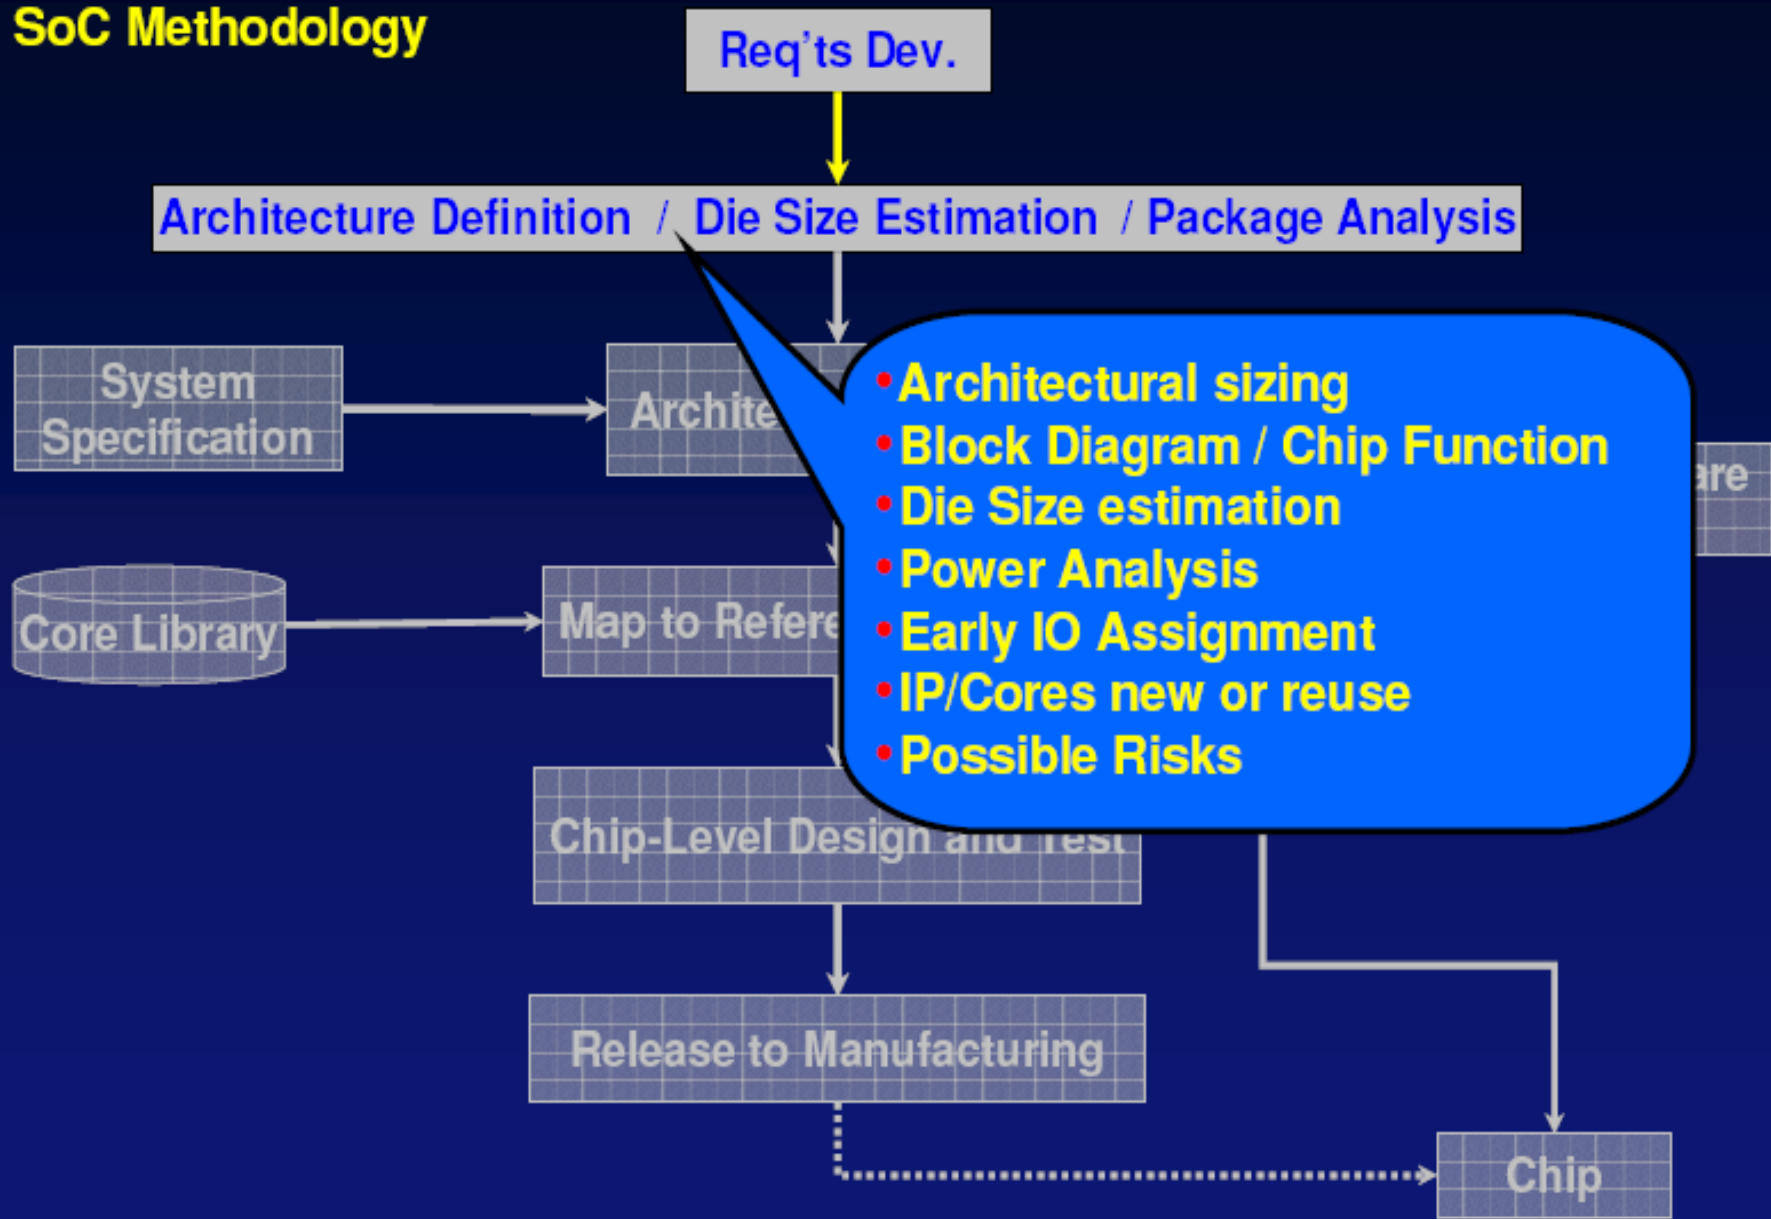
SOC Methodology Evolving ...

SoC Methodology



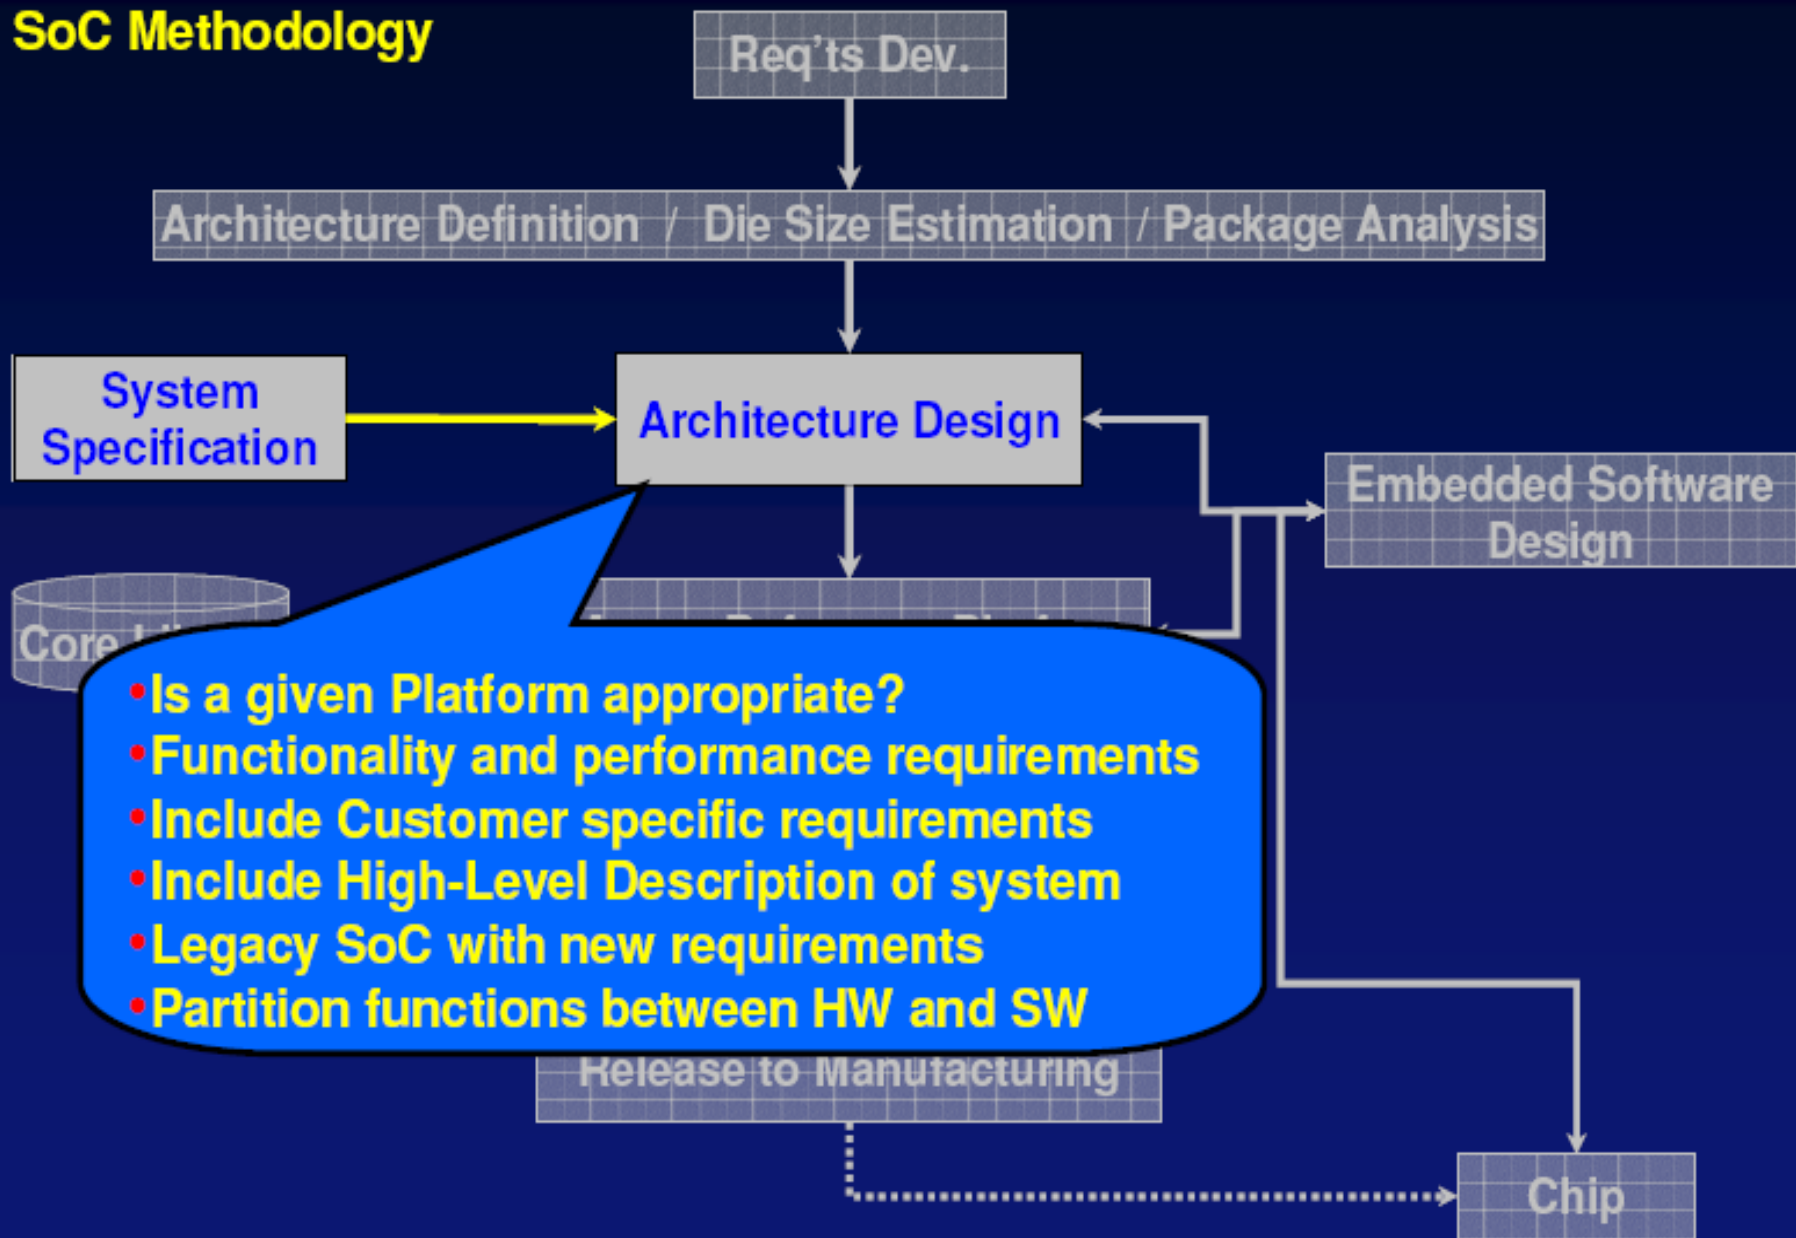
How to Design an SOC

SoC Methodology



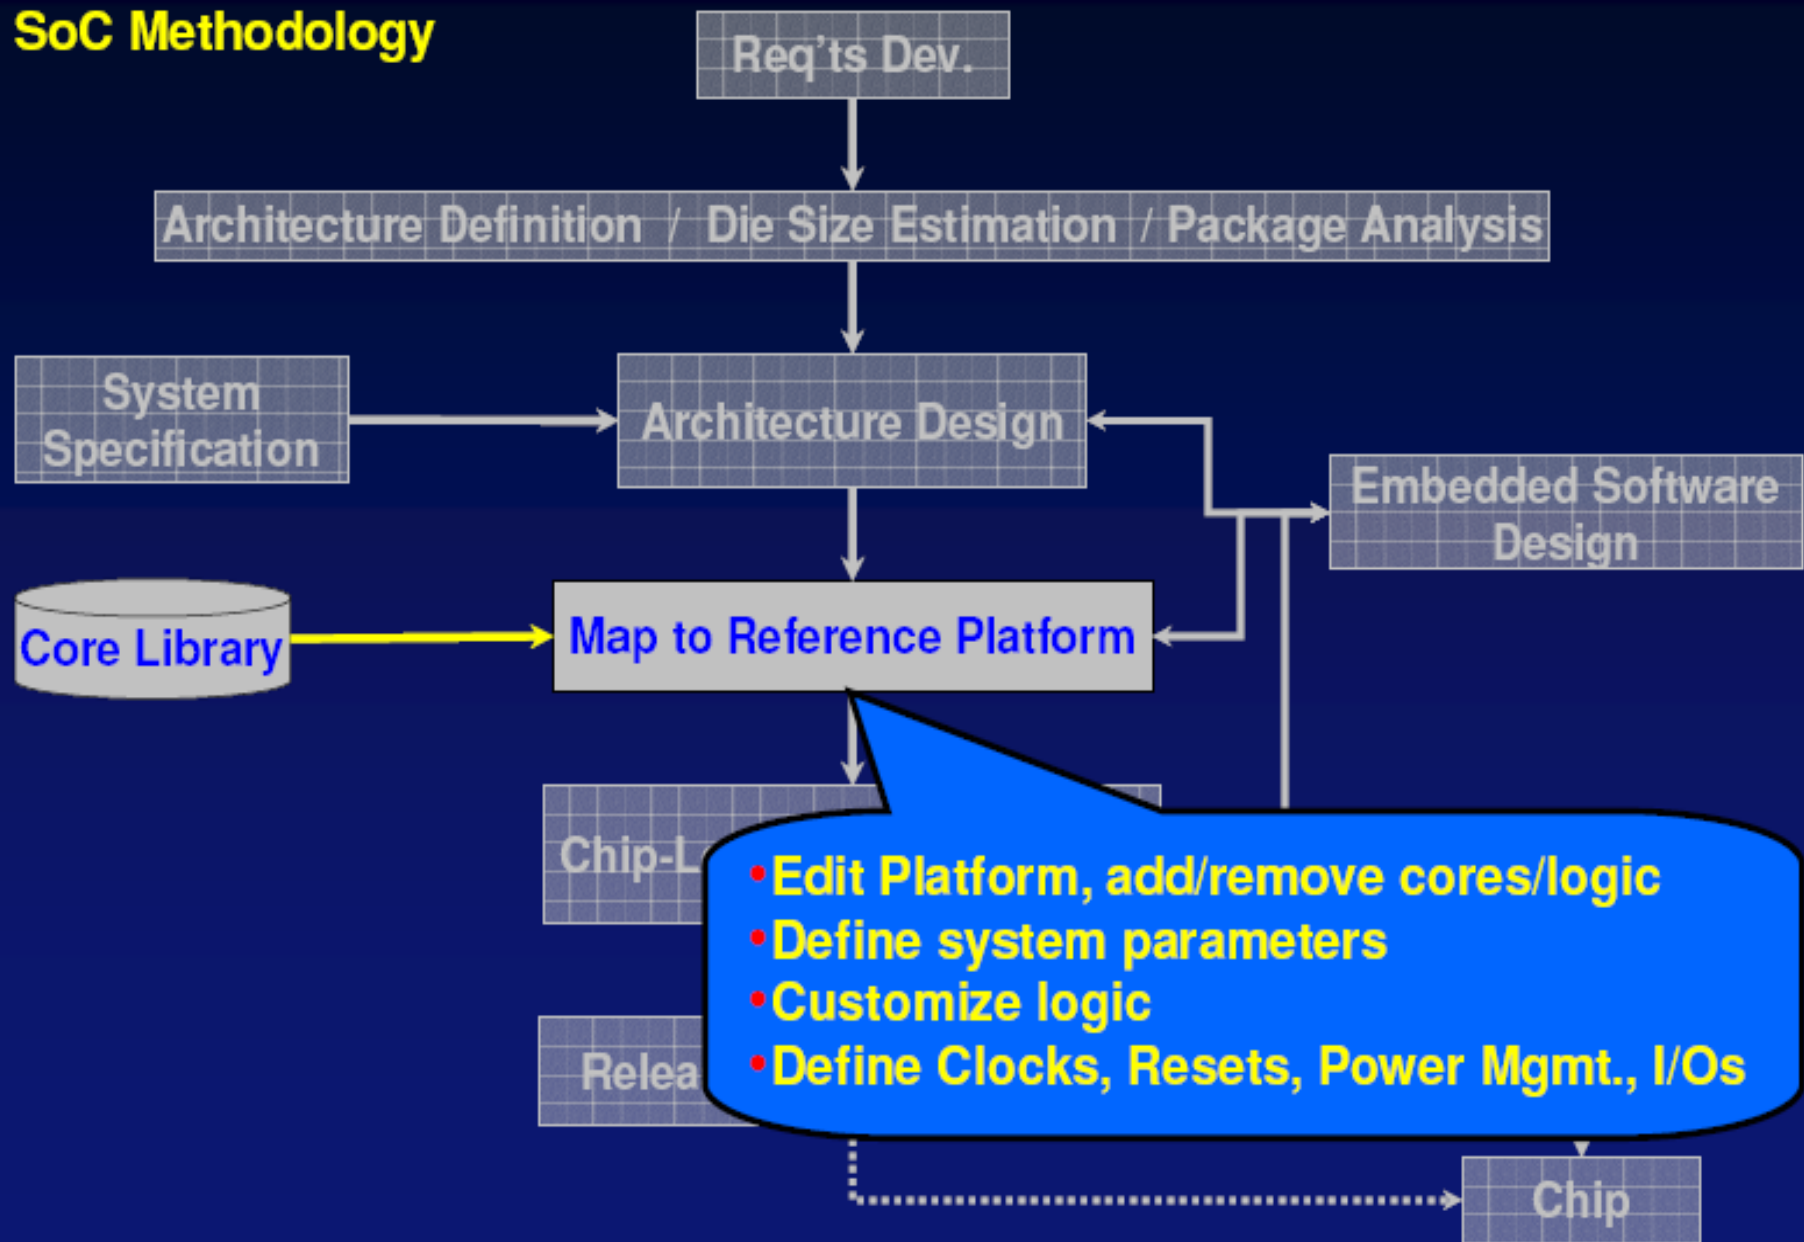
How to Design an SOC

SoC Methodology



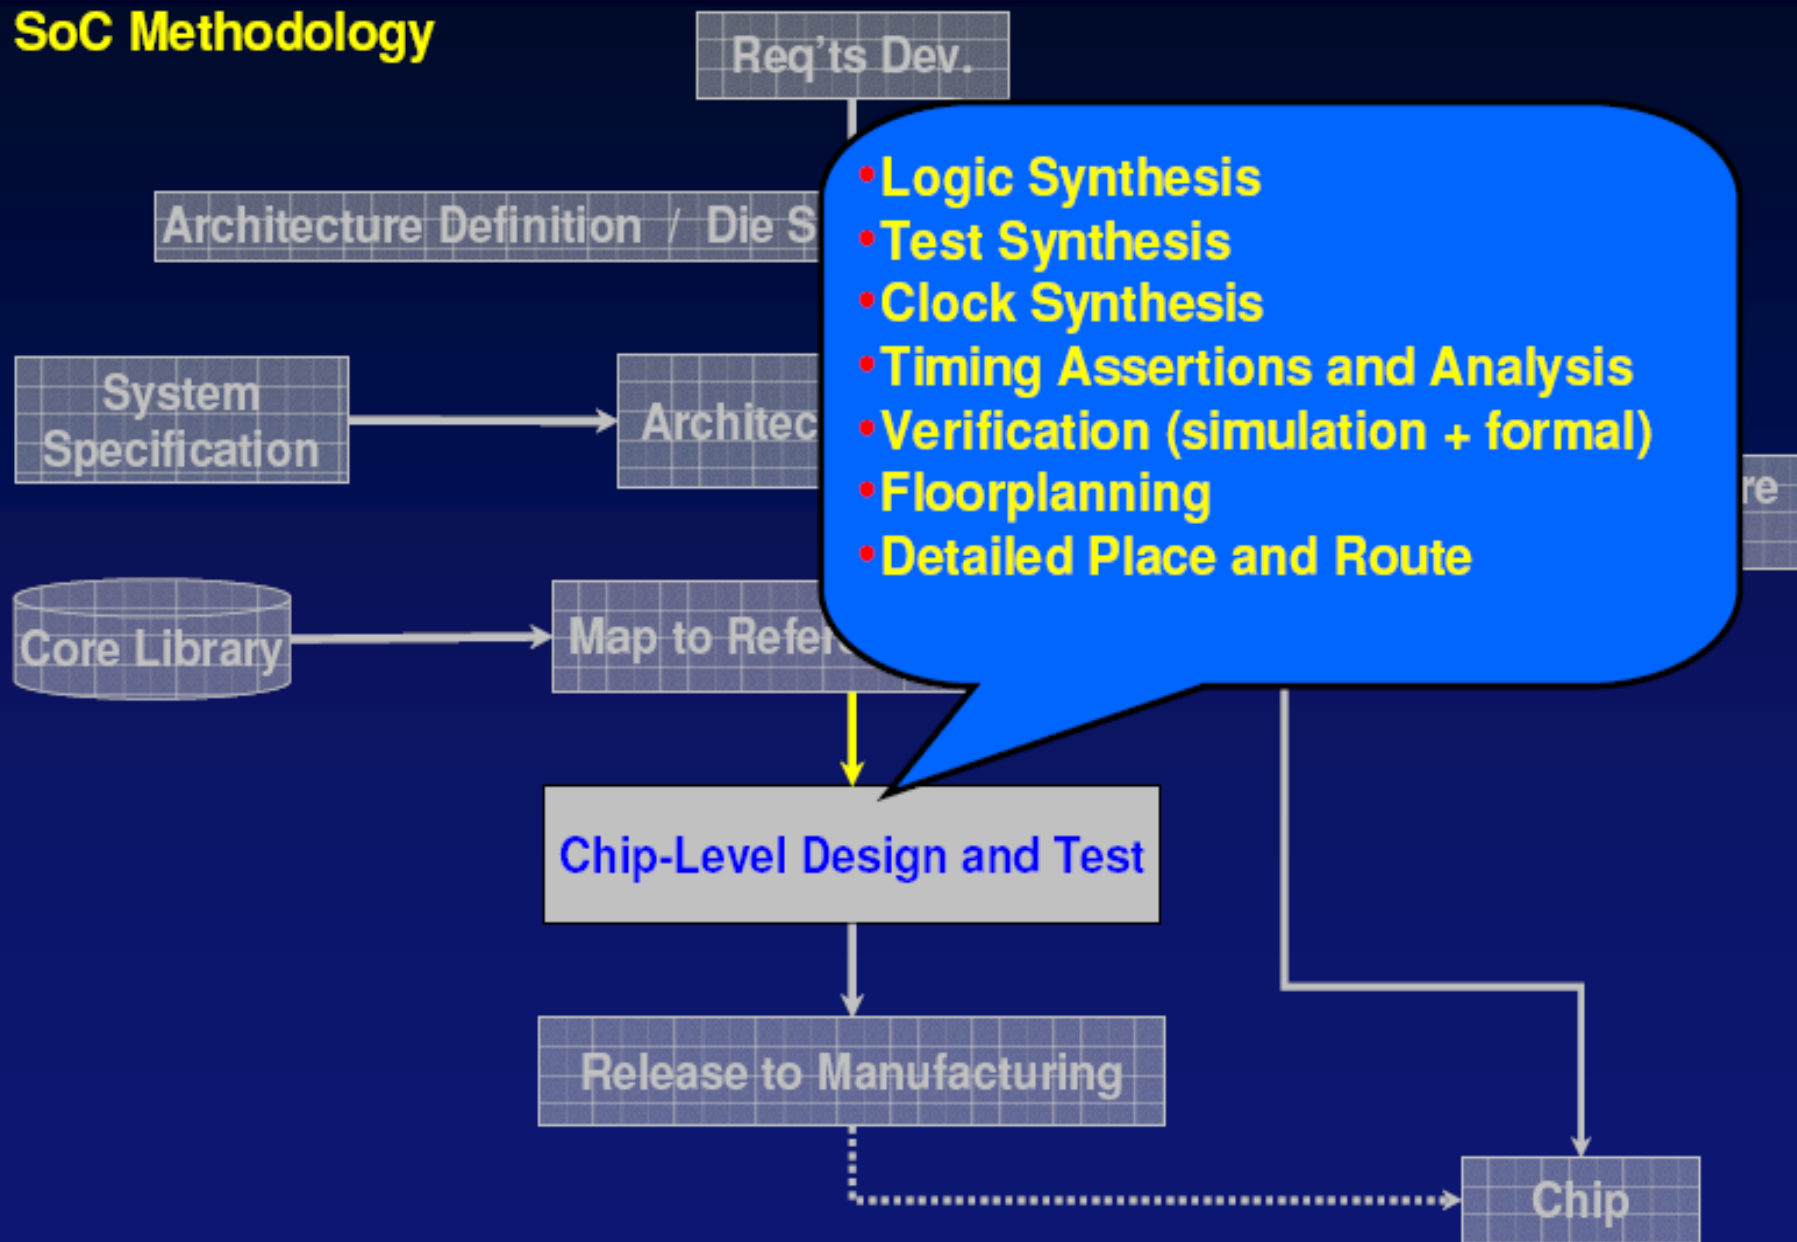
How to Design an SOC

SoC Methodology



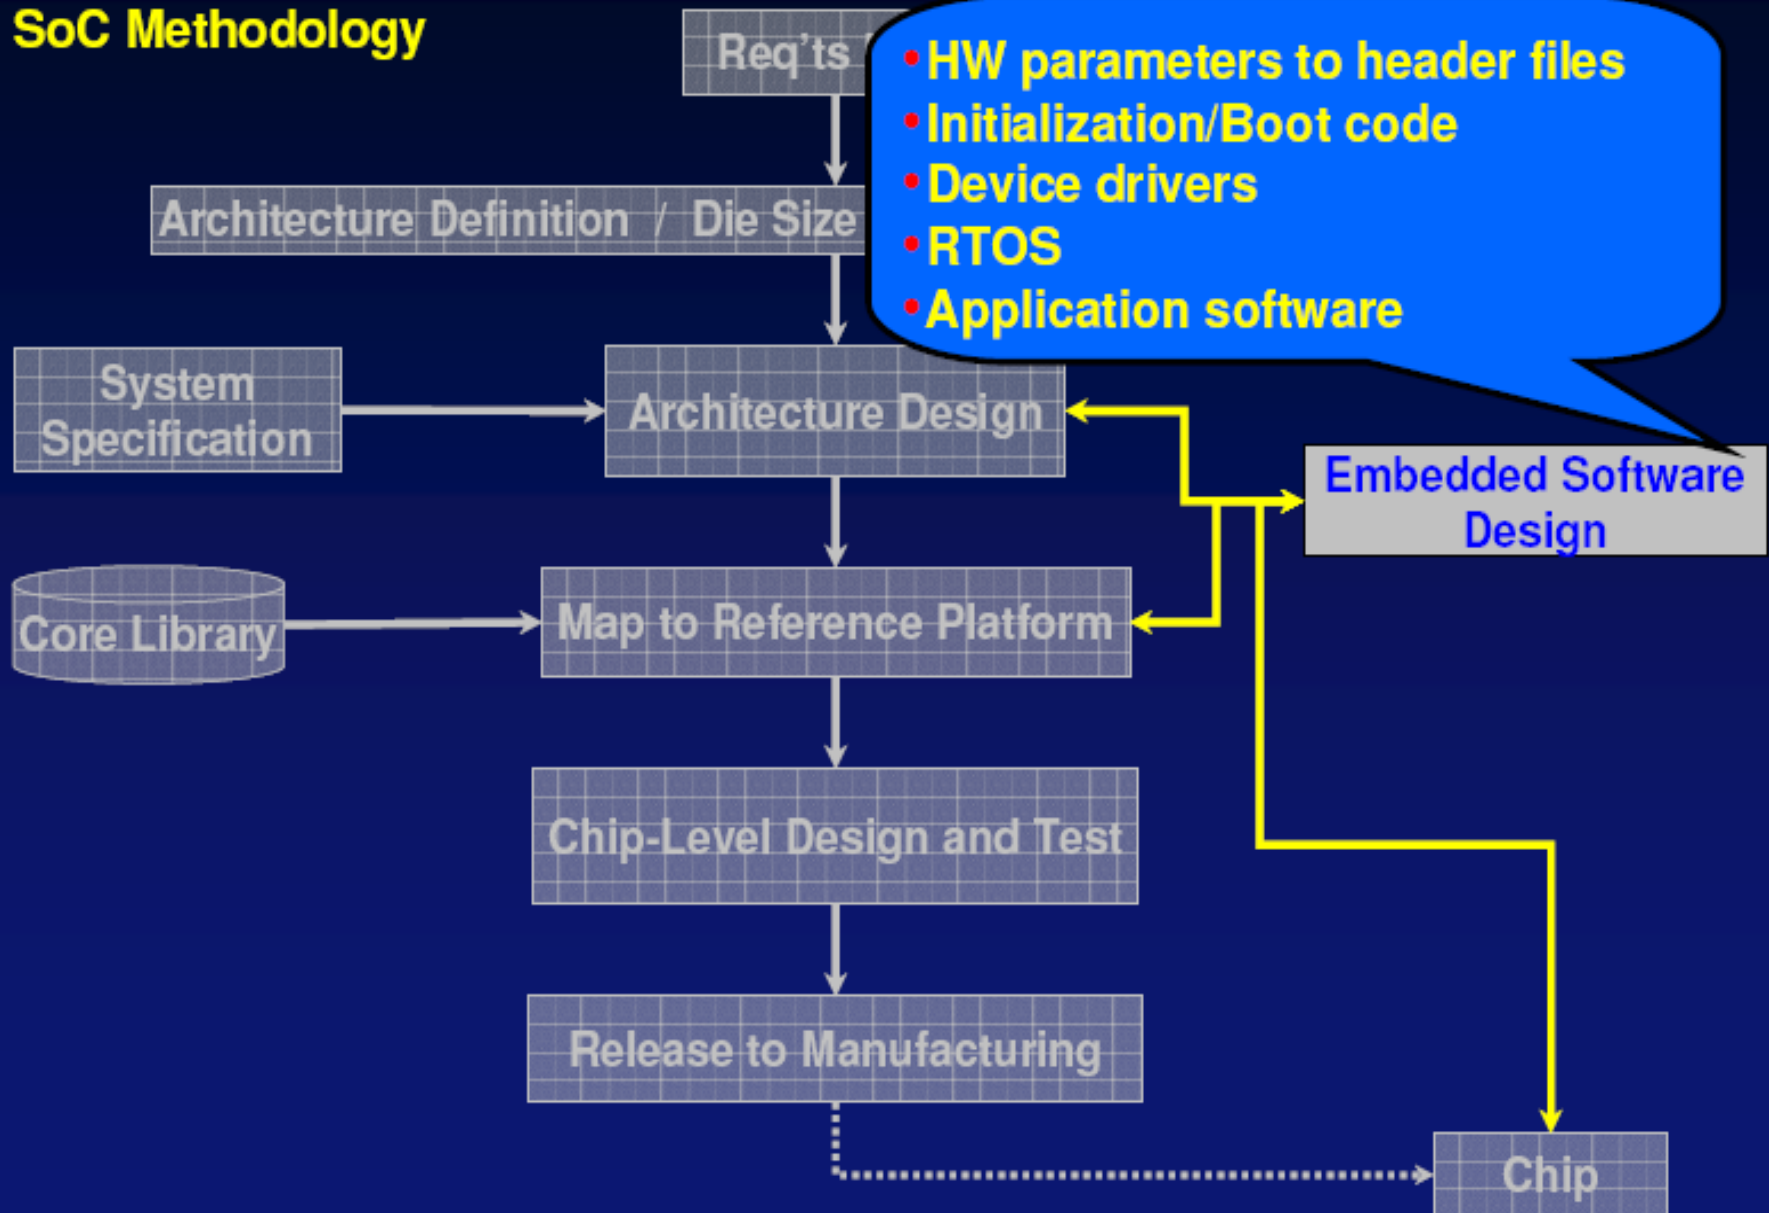
How to Design an SOC

SoC Methodology




How to Design an SOC

SoC Methodology



System on Chip - Testing

- SOCs are complex designs combining logic, memory and mixed-signal circuits in a single IC



Main SOC testing challenges

- Core level test: Embedded cores are tested as a part of the system
- Test access: Due to absence of physical access to the core peripherals, electronic access mechanism required
- SOC level test: SOC test is a single composite test including individual core, and UDL test and test scheduling

Test data volume for core-based SOC designs is very high.

- New techniques are required to reduce testing time, test cost, and the memory requirements of the automatic test equipment (ATE)

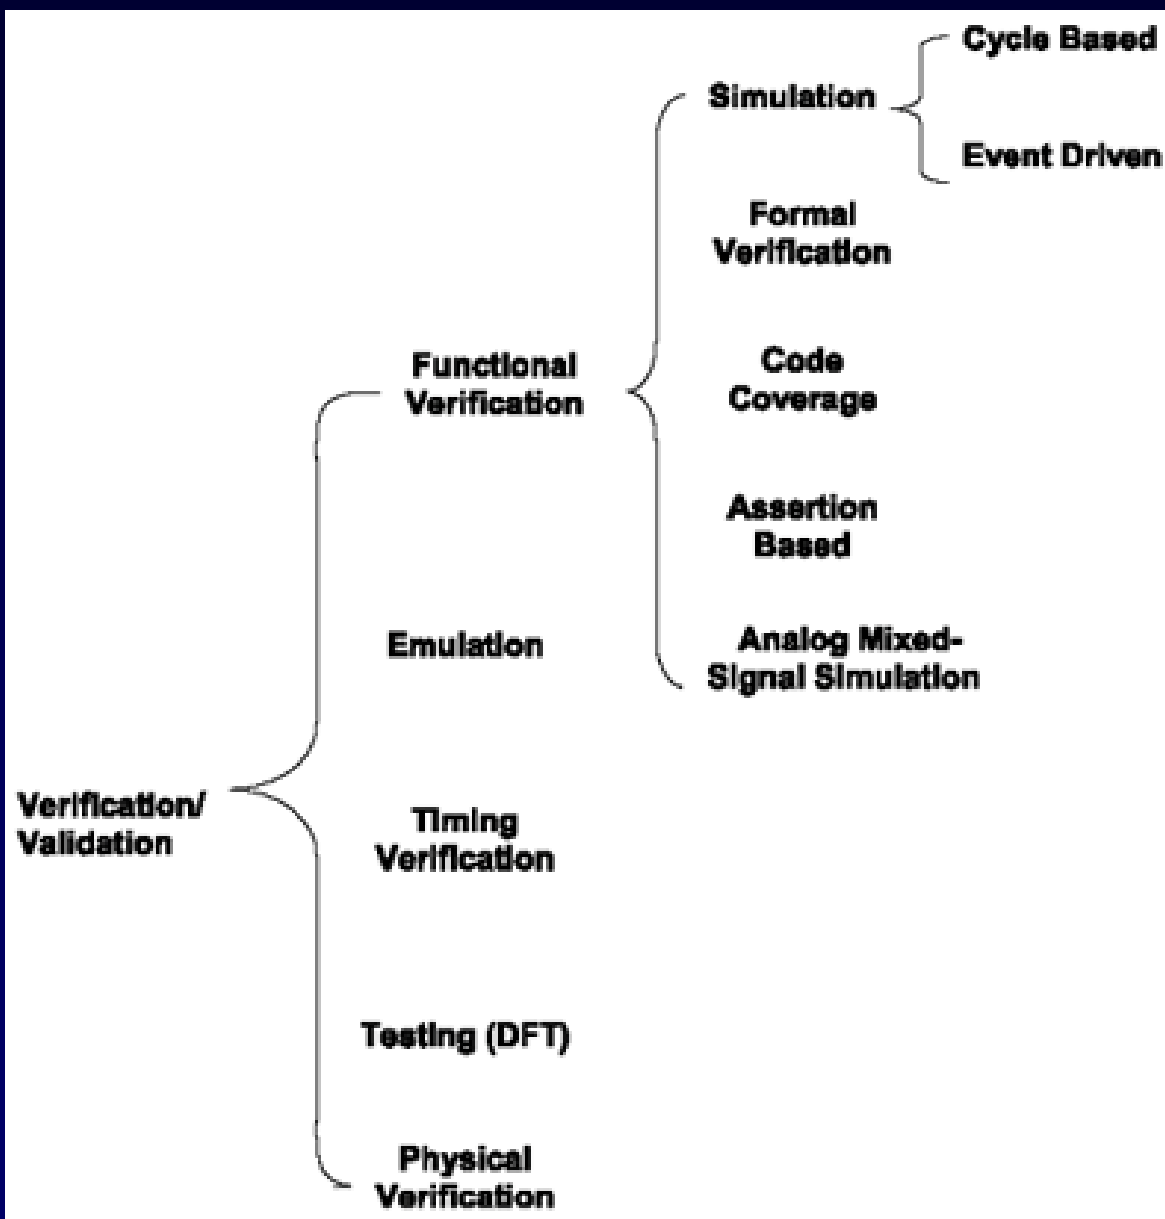
Verification

Today about 70% of design cost and effort is spent on verification.

Verification teams are often almost twice as large as the RTL designers at companies developing ICs.

Traditionally, chip design verification focuses on simulation.

However, new verification techniques are emerging.



Design for Integration

A key issue in SOC design is integration of silicon IPs (cores).

Integration of IPs directly affects the complexity of SOC designs and also influences verification of the SOC.

Verification is faster and easier if the SOC interconnect is simple and unified (use an on-chip communication system or intelligent on-chip bus).

There is no standard for OCBs; they are chosen almost exclusively by the specific application for which they will be used and by the designer's preference.

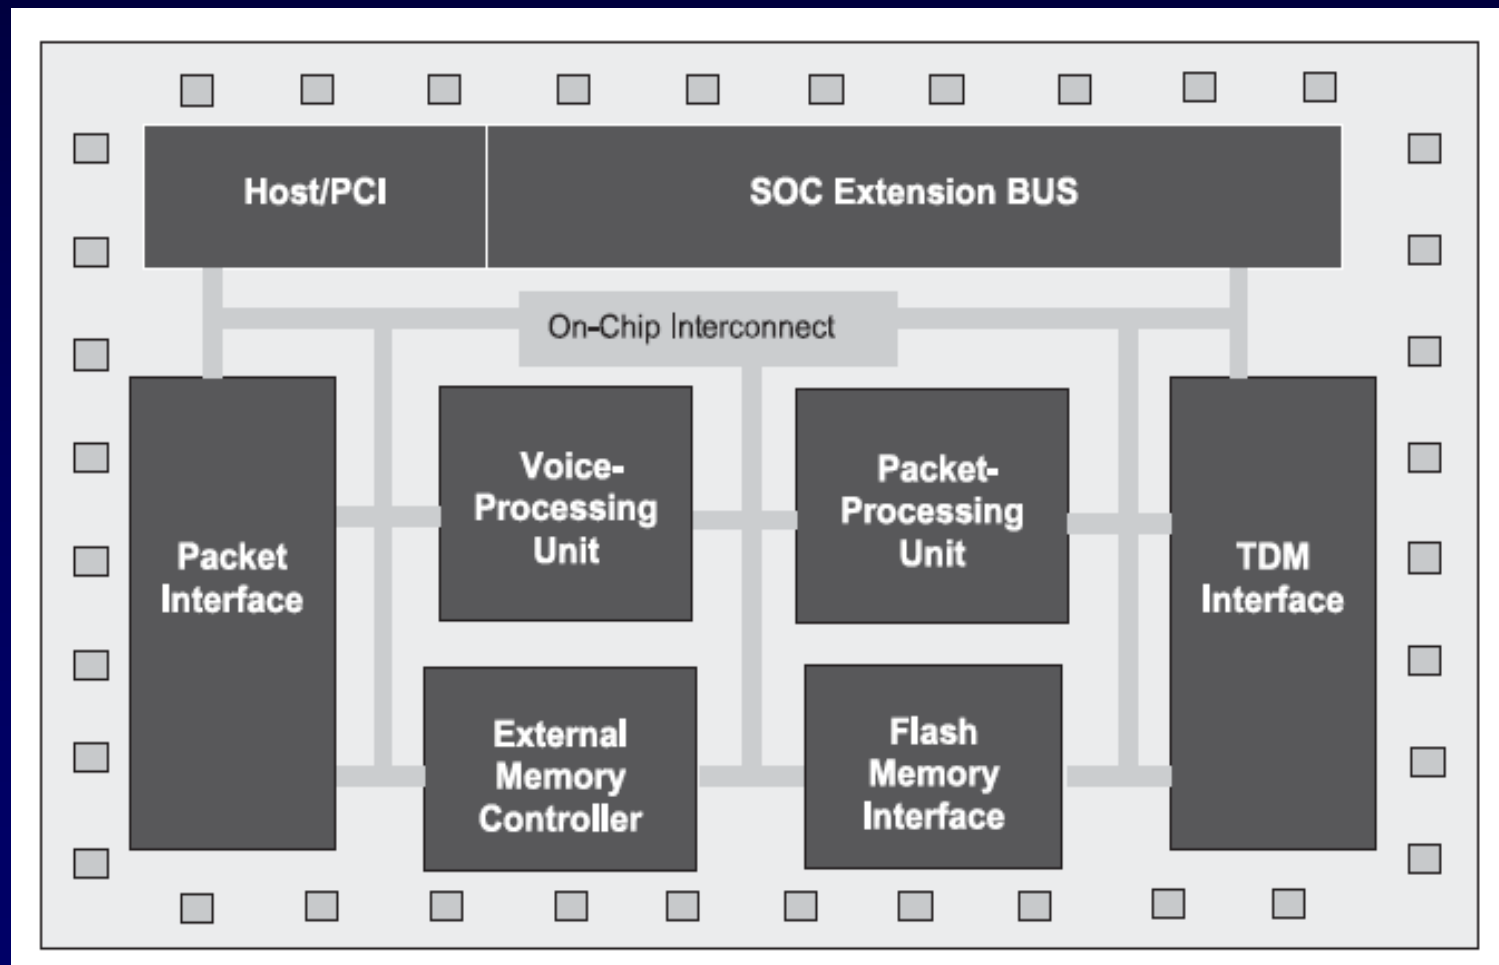
Two main types of OCBs (on-chip bus) and their characteristics

OCB	Speed	Bandwidth	Arbitration	Example
System	High	High	Complex	ARM AHB
Peripheral	Low	Low	Simple	PCI Bus

A Typical Gateway SoC Architecture

An example of typical gateway VoIP (Voice over Internet Protocol) system-on-a-chip diagram.

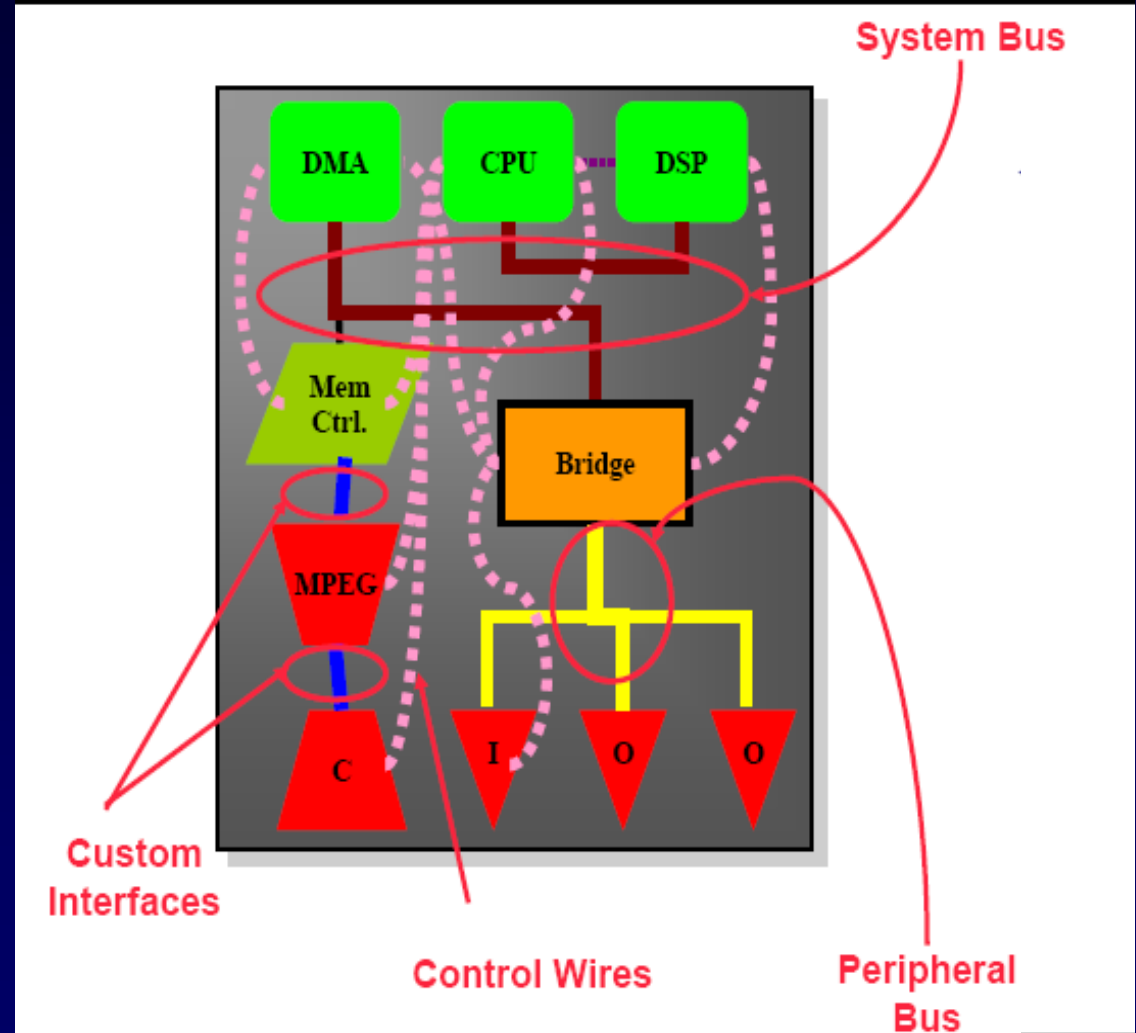
A gateway VoIP SoC is a device used for functions such as vocoders, echo cancellation, data/fax modems, and VoIP protocols.



A Traditional SOC Architecture (bus-based)

In a typical SOC, there are complex data flows and multiple cores such as CPUs, DSPs, DMA, and peripherals.

Therefore, resource sharing becomes an issue, communication between IPs becomes very complicated.



Sonics' SiliconBackplane Used in SOC Design Architecture

The CPU, DMA, and the DSP engine all share the same bus (the CPU or the system bus). Also, there are dedicated data links, a lot of control wires between blocks, and peripheral buses between subsystems

⇒ there is interdependency between blocks and a lot of wires in the chip.

Therefore, verification, test, and physical design all become difficult to fulfill.

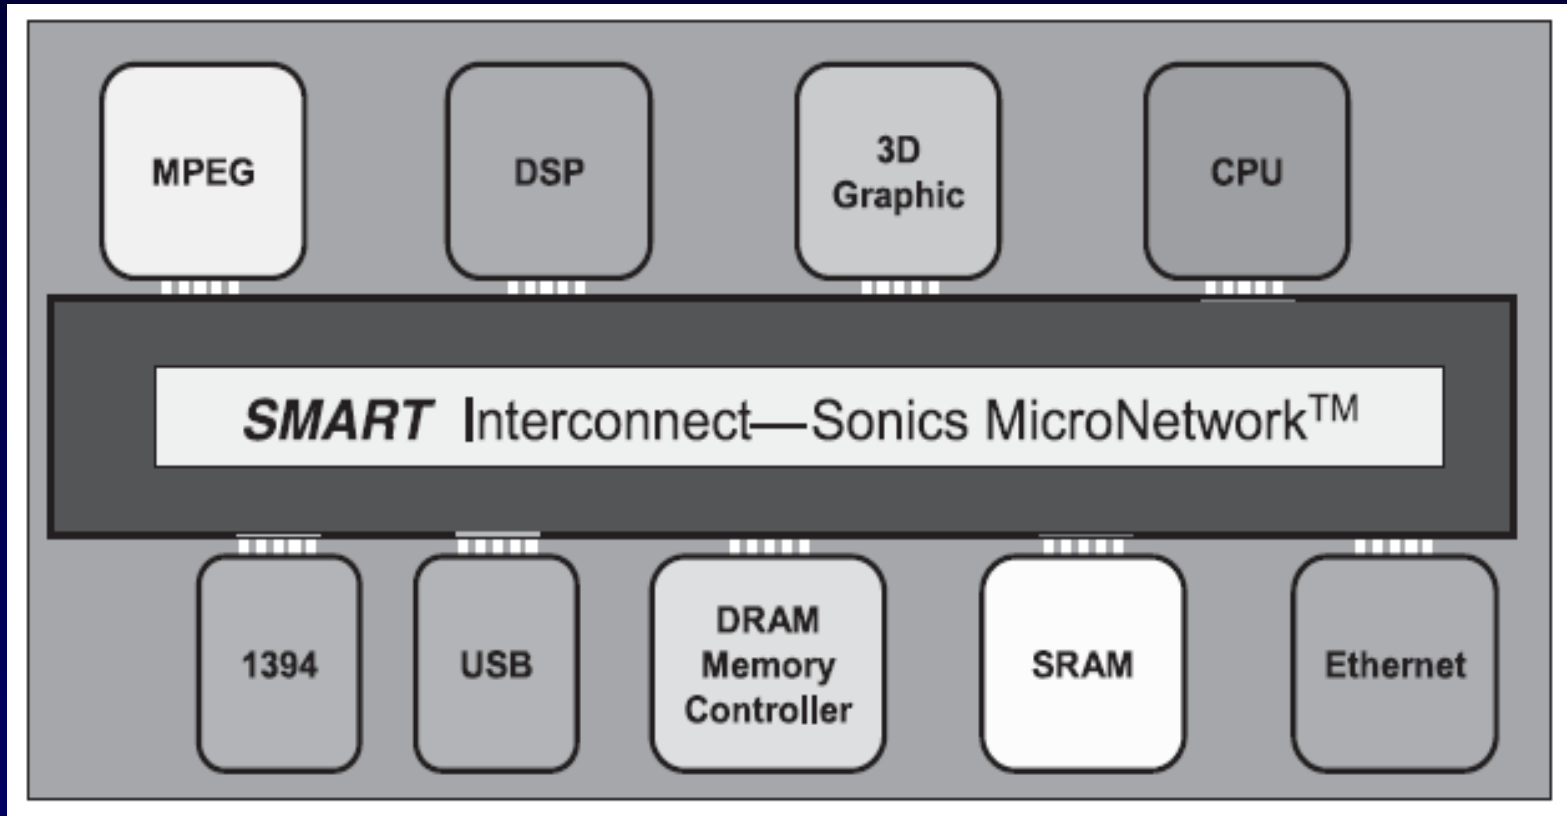
A solution to this system integration is to use an intelligent, on-chip interconnect that unifies all the traffic into a single entity.

An example of this is Sonics' SMART Interconnect SiliconBackplane MicroNetwork.

When compared to a traditional CPU bus, an on-chip interconnect such as Sonics SiliconBackplane has the following advantages:

- ❖ Higher efficiency
- ❖ Flexible configuration
- ❖ Guaranteed bandwidth and latency
- ❖ Integrated arbitration

Sonics' SiliconBackplane MicroNetwork Used in SOC Design Architecture


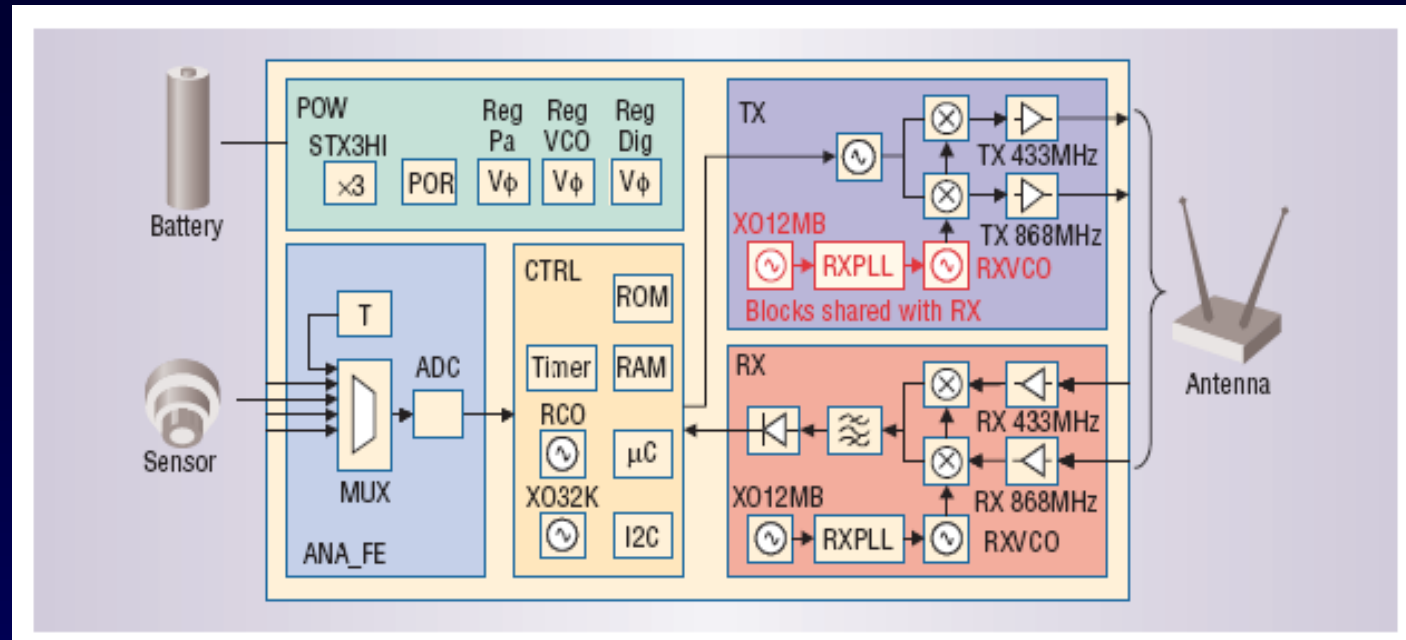


A MicroNetwork is a heterogeneous, integrated network that unifies, decouples, and manages all of the communication between processors, memories, and input/output devices.

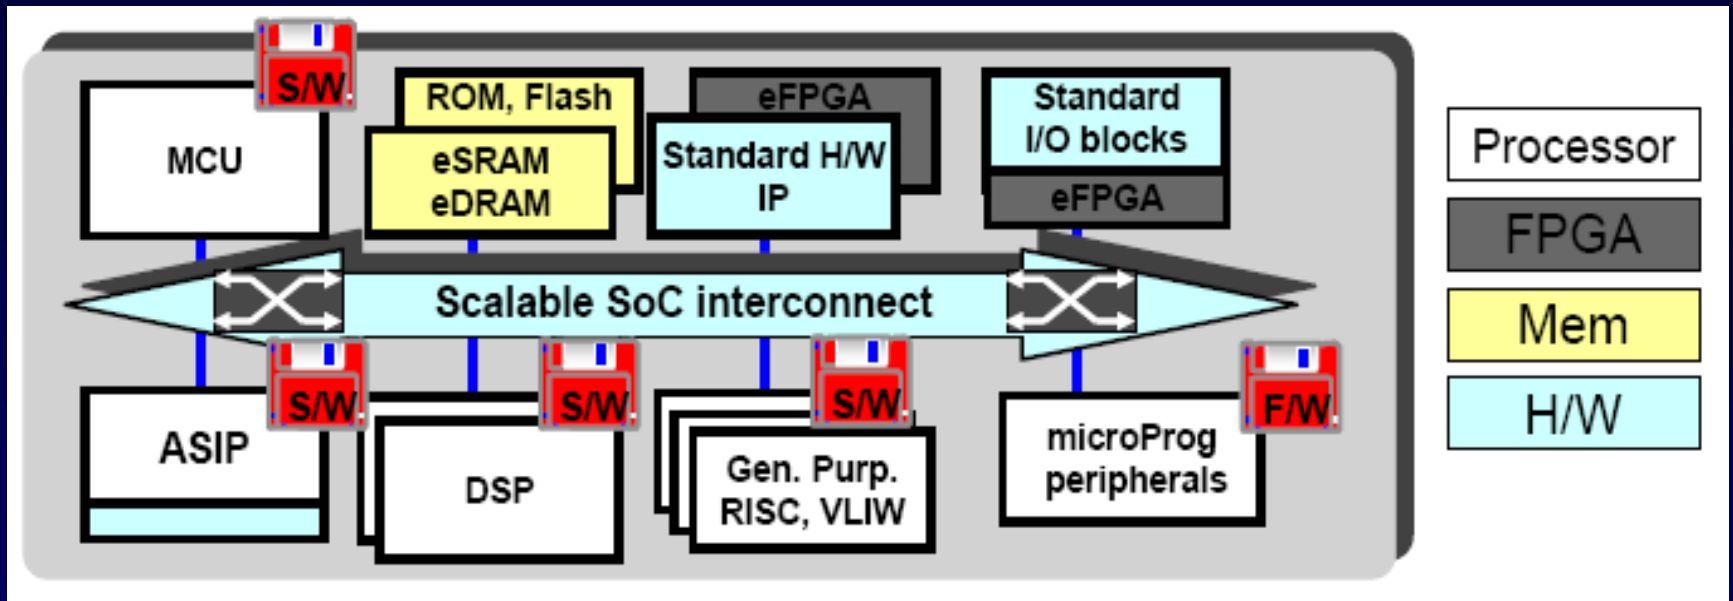
The basic WiseNET SoC architecture

The architecture includes:

- the ultralow-power dual-band radio transceiver (Tx and Rx),
- a sensor interface with a signal conditioner and two analog-to-digital converters (ANA_FE),
- a digital control unit based on a Cool-RISC microcontroller (μ C) with on-chip low-leakage memory, several timebasis and digital interfaces,
- a power management block (POW)



Networks on a chip



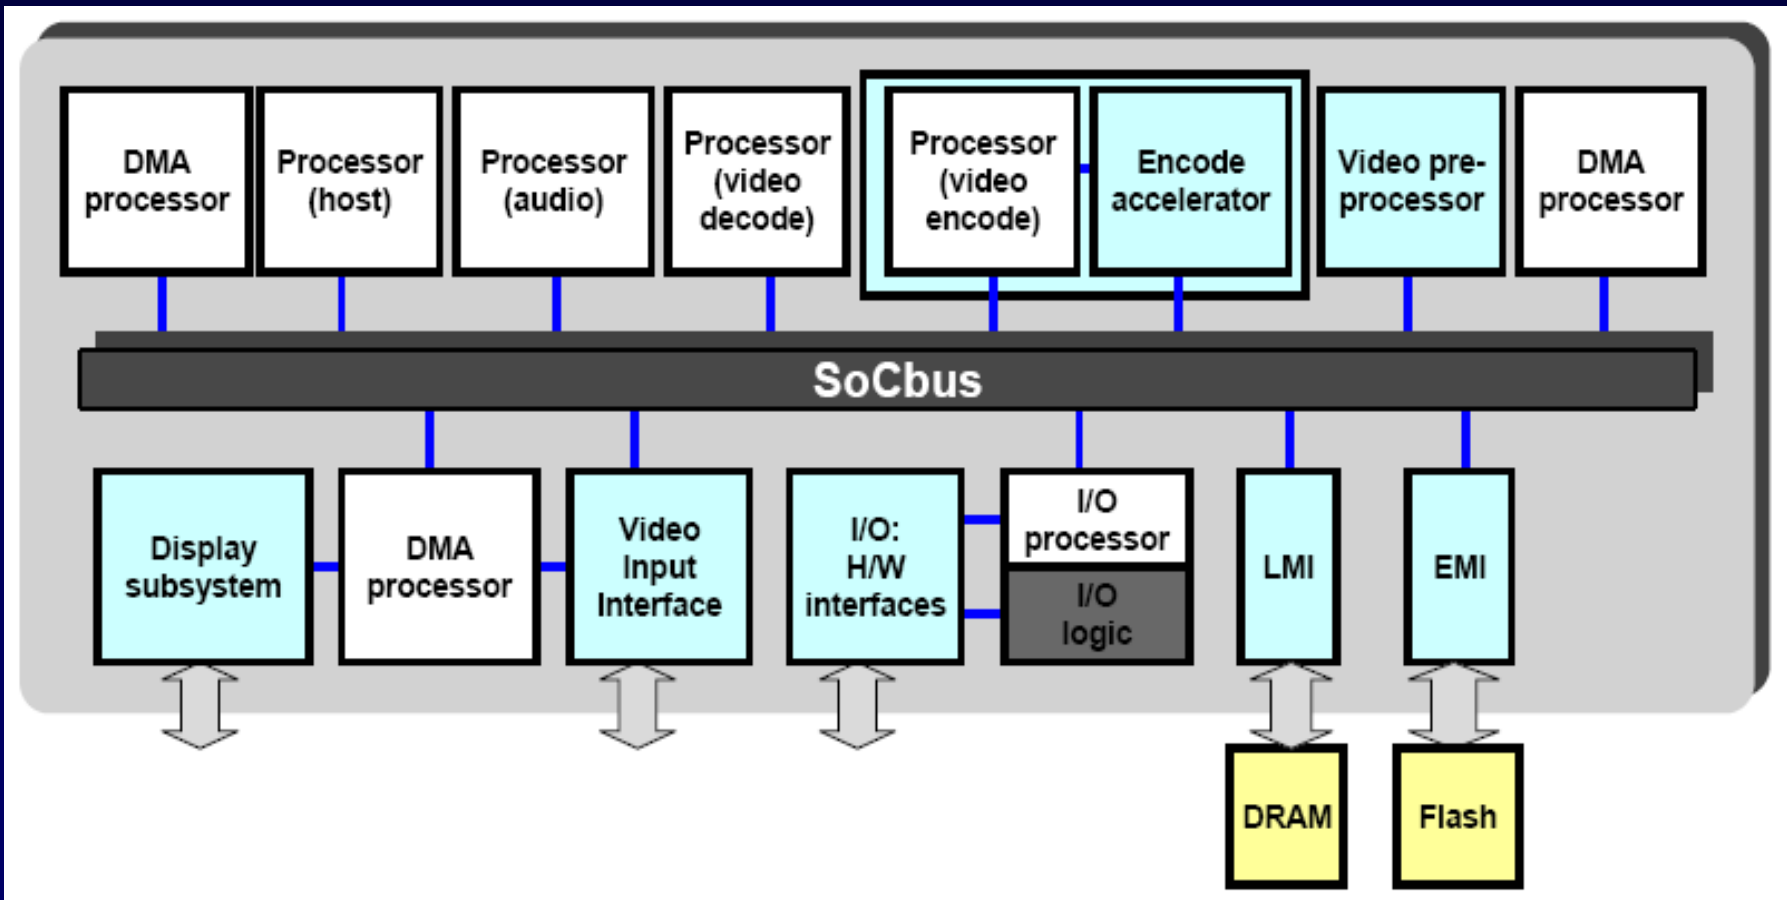
SoC for DVB

H/W

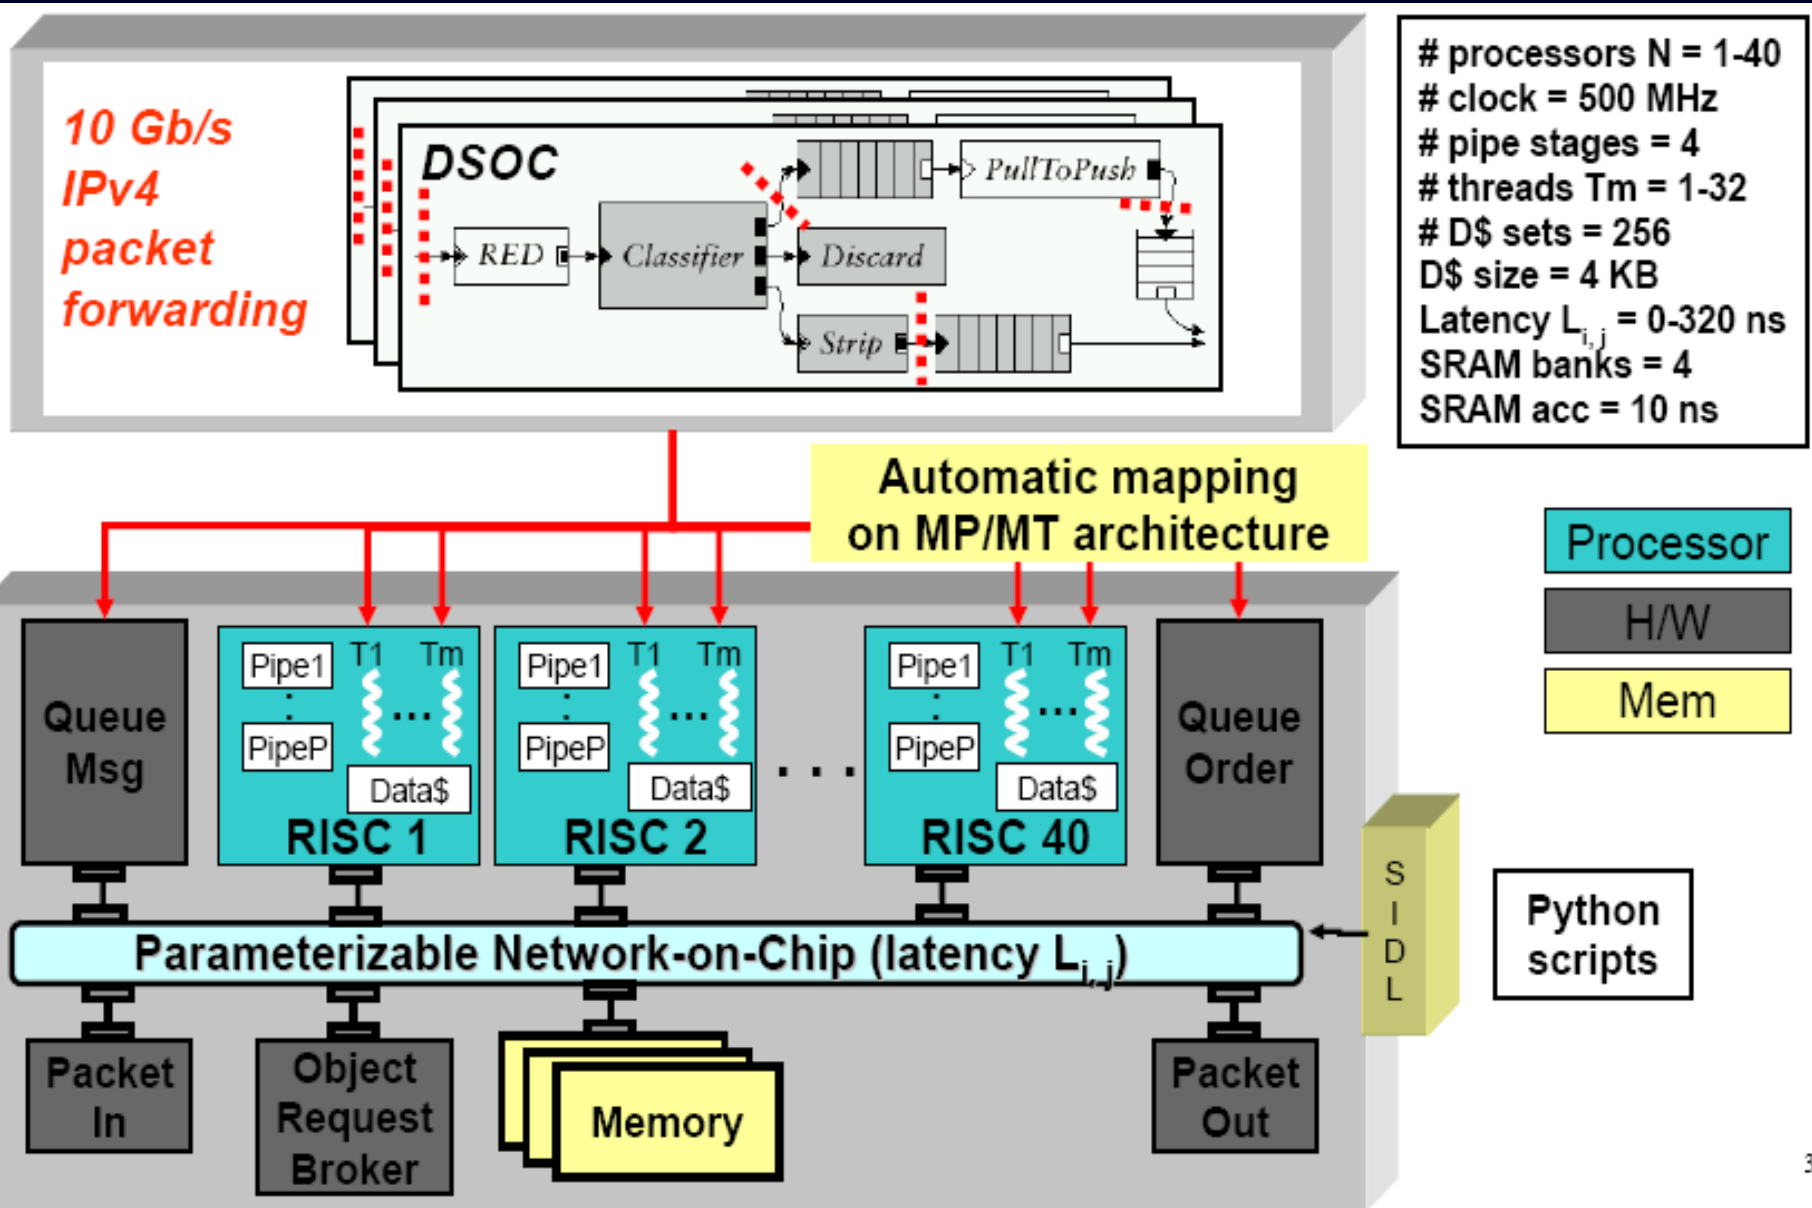
Processor

Interconnect

Soft logic



Network Processor



Conclusions

- An System on Chip (SoC) is an integrated circuit that implements most or all of the function of a complete electronic system.

Four vital areas of SoC:

- ❖ Higher levels of abstraction
- ❖ IP and platform re-use
- ❖ IP creation – ASIPs, interconnect and algorithm
- ❖ Earlier software development and integration

